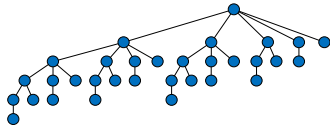


Lecture 3. Priority Queues

CpSc 8400: Algorithms and Data Structures
Brian C. Dean

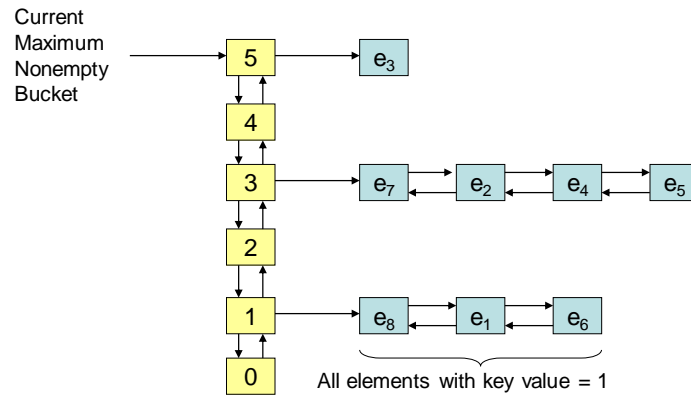


School of Computing
Clemson University
Spring, 2016

Warm-Up: Incremental Priority Queues

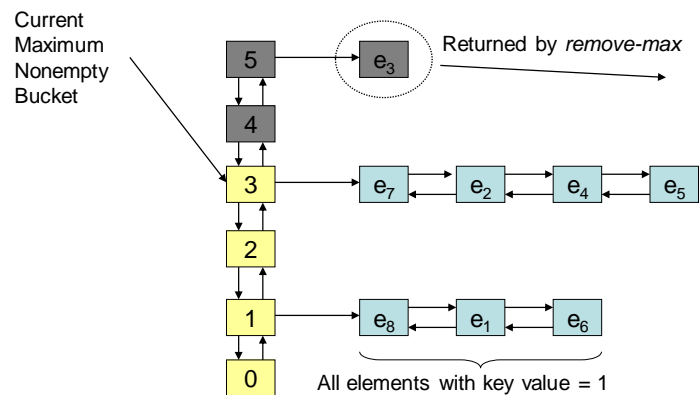
- Fundamental operations of a (max-) priority queue:
 - *Insert* : insert new element
 - *Remove-max* : remove element with maximum key
- We'll study general priority queues in a moment, but for now, consider the special case of an *incremental* priority queue:
 - Keys stored in the structure are nonnegative integers, initially zero.
 - We additionally support an *increment-priority(e)* operation that takes a pointer to an element and increases its key by 1.

Implementing an Incremental Priority Queue



3

The Remove-Max Operation



4

Analysis of Incremental Priority Queue

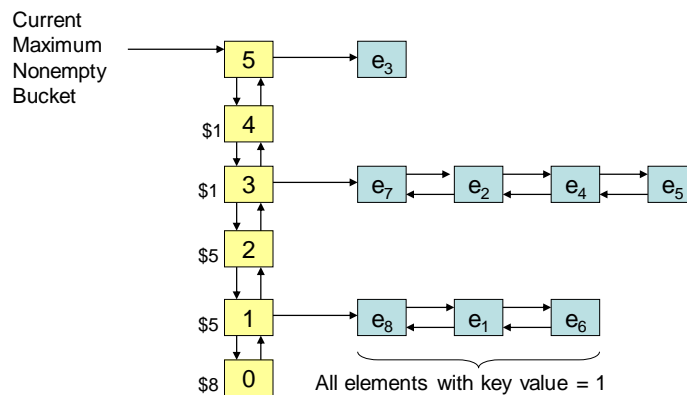
- Let M denote the amount by which the “current maximum bucket” pointer moves.

	Worst-Case Running Time	Amortized Running Time
<i>insert</i>	1	1
<i>increment-priority</i>	1	2
<i>remove-max</i>	$1+M$ (not bounded!)	1

↑
All operations have $O(1)$ amortized running times!

5

Plenty of Credit to Spare...



6

General Priority Queues

- In a simple FIFO queue, elements exit in the same order as they enter.
- In a priority queue, the element with highest priority (usually defined as having *lowest* key) is always the first to exit.
- Many uses:
 - **Scheduling:** Manage a set of tasks, where you always perform the highest-priority task next.
 - **Sorting:** Insert n elements into a priority queue and they will emerge in sorted order.
 - **Complex Algorithms:** For example, Dijkstra's shortest path algorithm is built on top of a priority queue.

7

Priority Queues

- All priority queues support:
 - Insert(e, k)* : Insert a new element e with key k .
 - Remove-Min* : Remove and return the element with minimum key.
- In practice (mostly due to Dijkstra's algorithm), many support:
 - Decrease-Key($e, \Delta k$)* : Given a pointer to element e within the heap, reduce e 's key by Δk .
- Some priority queues also support:
 - Increase-key($e, \Delta k$)* : Increase e 's key by Δk .
 - Delete(e)* : Remove e from the structure.
 - Find-min* : Return a pointer to the element with minimum key.

8

Redundancies Among Operations

- Given *insert* and *delete*, we can implement *increase-key* and *decrease-key*.
- Given *decrease-key* and *remove-min*, we can implement *delete*.
- Given *find-min* and *delete*, we can implement *remove-min*.
- Given *insert* and *remove-min*, we can implement *find-min*.

9

Priority Queue Implementations



- There are *many* simple ways to implement the abstract notion of a priority queue as a concrete data structure:

	<i>insert</i>	<i>remove-min</i>
Unsorted array or linked list	$O(1)$	$O(n)$
Sorted array or linked list	$O(n)$	$O(1)$
Binary heap	$O(\log n)$	$O(\log n)$
Balanced binary search tree	$O(\log n)$	$O(\log n)$
Skew heap	$O(\log n)$ am.	$O(\log n)$ am.

10

The Binary Heap

- An almost-complete binary tree (all levels full except the last, which is filled from the left side up to some point).
- Satisfies the **heap property**: for every element e , $\text{key}(\text{parent}(e)) \leq \text{key}(e)$.
 - Minimum element always resides at root.
- Physically stored in an array $A[0 \dots n-1]$.
- Easy to move around the array in a treelike fashion:
 - $\text{Parent}(i) = \text{floor}((i-1)/2)$.
 - $\text{Left-child}(i) = 2i + 1$
 - $\text{Right-child}(i) = 2i + 2$.



11

Heap Operations : sift-up and sift-down

- All binary heap operations are built from the two fundamental operations *sift-up* and *sift-down*:
 - *sift-up*(i) : Repeatedly swap element $A[i]$ with its parent as long as $A[i]$ violates the heap property with respect to its parent (i.e., as long as $A[i] < A[\text{parent}(i)]$).
 - *sift-down*(i) : As long as $A[i]$ violates the heap property with one of its children, swap $A[i]$ with its smallest child.
- Both operations run in $O(\log n)$ time since the height of an n -element heap is $O(\log n)$.
- In some other places, *sift-down* is called *heapify*, and *sift-up* is known as *up-heap*.


12

Implementing Heap Operations Using sift-up and sift-down

- The remaining operations are now easy to implement in terms of *sift-up* and *sift-down*:
 - *insert* : place new element in $A[n+1]$, then *sift-up*($n+1$).
 - *remove-min* : swap $A[n]$ and $A[1]$, then *sift-down*(1).
 - *decrease-key*($i, \Delta k$) : decrease $A[i]$ by Δk , then *sift-up*(i).
 - *increase-key*($i, \Delta k$) : increase $A[i]$ by Δk , then *sift-down*(i).
 - *delete*(i) : swap $A[i]$ with $A[n]$, then *sift-up*(i), *sift-down*(i).
- All of these clearly run in $O(\log n)$ time.
- General idea: modify the heap, then fix any violation of the heap property with one or two calls to *sift-up* or *sift-down*.

13

Caveat: You Can't Easily Find Elements In Heaps (Except the Min)



Each record in the data structure keeps a pointer to the physical element of data it represents, and each element of data maintains a pointer to its corresponding record in the data structure.

14


Building a Binary Heap

- We could build a binary heap in $O(n \log n)$ time using n successive calls to *insert*.
- Another way to build a heap: start with our n elements in arbitrary order in $A[0..n-1]$, then call *sift-down*(i) for $i = n-1$ down to 0.
 - Remarkable fact #1: this builds a valid heap!
 - Remarkable fact #2: this runs in only $O(n)$ time!

15

Bottom-Up Heap Construction

- The key property of *sift-down* is that it fixes an isolated violation of the heap property at the root:



- Using induction, it is now easy to prove that our “bottom-up” construction yields a valid heap.

16

Bottom-Up Heap Construction

- To analyze the running time of bottom-up construction, note that:
 - At most n elements reside in the bottom level of the heap. Only 1 unit of work done to them by *sift-down*.
 - At most $n/2$ elements reside in the 2nd lowest level, and at most 2 units of work are done to each of them.
 - At most $n/4$ elements reside in the 3rd lowest level, and at most 3 units of work are done to them.
- So total time $\leq T = n + 2(n/2) + 3(n/4) + 4(n/8) + \dots$
(for simplicity, we carry the sum out to infinity, as this will certainly give us an upper bound).
- Claim: $T = 4n = O(n)$

17

“Shifting” Technique for Sums

$$\begin{array}{rcl}
 T & = & n + 2(n/2) + 3(n/4) + 4(n/8) + \dots \\
 - \quad T/2 & = & \quad n/2 + 2(n/4) + 3(n/8) + \dots \\
 \hline
 T/2 & = & n + n/2 + n/4 + n/8 + \dots
 \end{array}$$

Applying the same trick again:

$$\begin{array}{rcl}
 T & = & 2n + n + (n/2) + (n/4) + \dots \\
 - \quad T/2 & = & \quad n + (n/2) + (n/4) + \dots \\
 \hline
 T/2 & = & 2n
 \end{array}$$

18


Heapsort

- Any priority queue can be used to sort. Just use n *inserts* followed by n *remove-mins*.
- The binary heap gives us a particularly nice way to sort in $O(n \log n)$ time, known as **heapsort**:
 - Start with an array $A[0..n-1]$ of elements to sort.
 - Build a heap (bottom up) on A in $O(n)$ time.
 - Call *remove-min* n times.
 - Afterwards, A will end up reverse-sorted (it would be forward-sorted if we had started with a “max” heap)

19

Another Simple Way to Implement Priority Queues...

- Suppose we store our priority queue in a “heap-ordered” binary tree.
 - Heap property: $\text{parent} \leq \text{child}$.
 - Each node maintains a pointer to its left child and right child.
 - The tree is not necessarily “balanced”. It could conceivably be nothing more than a single sorted path.
 - No longer easily mapped to an array, as with a binary heap.



20


All You Need is Merge...

- Suppose we can *merge* two heap-ordered trees in $O(\log n)$ time.
- All priority queue operations now easy to implement in $O(\log n)$ time!
 - *insert*: merge with a new 1-element tree.
 - *remove-min*: remove root, merge left & right subtrees.

21

Merging Two Heap-Ordered Trees

- Take two heap-ordered trees h_1 and h_2 , where h_1 has the smaller root.
- Clearly, h_1 's root must become the root of the merged tree.
- To complete the merge, recursively merge h_2 into either the left or right subtree of h_1 :




- As a base case, the process ends when we merge a heap h_1 with an empty heap, the result being just h_1 .

22

Skew Heaps

- Always merge into R, but after merging h_2 into h_1 , just swap h_1 's children (so we really alternate between L and R).
- Remarkably, this makes merge (and therefore all other major operations) run in just $O(\log n)$ amortized time!



23

What About Delete / Decrease-Key?

- In a skew heap, *insert* and *remove-min* are based on merging, so they run in $O(\log n)$ amortized time.
- For *decrease-key* (and *increase-key*), simply delete an element and re-insert it with a new key.
- Subtle question: how do we implement *delete* efficiently? (so that it doesn't interfere with the amortized analysis of other operations...)

24