


Lecture 4. Sorting, Models of Computation, Lower Bounds

CpSc 8400: Algorithms and Data Structures
Brian C. Dean



School of Computing
Clemson University
Spring, 2016

Models of Computation

- **Question:** What is the worst-case running time of the fastest possible algorithm for sorting n numbers?
- A. $O(n \log n)$
B. $O(n (\log \log n)^{1/2})$
C. $O(n)$
D. $O(1)$

Models of Computation

- **Answer:** It depends on our model of computation!
- A. In the comparison model: $O(n \log n)$
- B. In the RAM model: $O(n (\log \log n)^{1/2})$ in expectation
- C. In the RAM model with small word size: $O(n)$
- D. In a (very unrealistic!) model where our machine has a “sort n integers” instruction: $O(1)$

3

Models of Computation

- In order to speak of the “running time” (number of fundamental operations) of an algorithm, we need to know what constitutes a fundamental operation.
- A *model of computation* defines the primitive operations in the abstract computing environment that our algorithm can use.
- Some algorithms look faster merely because they assume a more powerful model of computation.
 - Example: the *element uniqueness* problem (determining if an n -element array contains n distinct elements) requires $\Omega(n \log n)$ time in the comparison model, but can be solved in $O(n)$ expected time on a RAM.
 - The $O(n)$ RAM algorithm is not necessarily “better”...

4

The Random Access Machine (RAM) Model

- Our default model of computation.
- Memory is a long array of *words* (b-bit integers).
- Can perform random access into memory: setting and retrieving a word based on its index takes $O(1)$ time.
- Simple arithmetic operations also take $O(1)$ time: adding two words, multiplying two words, comparing two words, etc...
- Fairly good abstract model of a modern digital computer. Or is it...?
- Tricky subtlety: what can we assume about b?

5


Other Models

- The **real RAM**.
 - Words can, if desired, hold real numbers. We can perform arithmetic on two real numbers in $O(1)$ time.
 - Not realistic, but a useful simple model for problems that might require real numbers.
 - e.g., geometric problems where irrational distances might arise
 - Roundoff errors must be considered when actually implementing a real RAM algorithm.
- The **comparison-based** model.
 - Input elements not necessarily numeric, so can't use mathematical operations on them. Two input elements can only be compared.
 - Nice model for designing sorting/searching algorithms.

6

Exotic “Models of Computation”



- Physical models
(e.g., ball + string network models for computing shortest paths)
- Chemical models
(e.g., DNA computing for Hamiltonian cycles)
- Biological models
(e.g., computing shortest paths with slime molds)
- Optical models
(e.g., computing a Fourier transform by diffraction of light)
- Quantum computation



Hecht. Optics, 3rd edition



L.M. Adleman. Molecular computation of solutions to combinatorial problems. *Science* 11 (1994).



7

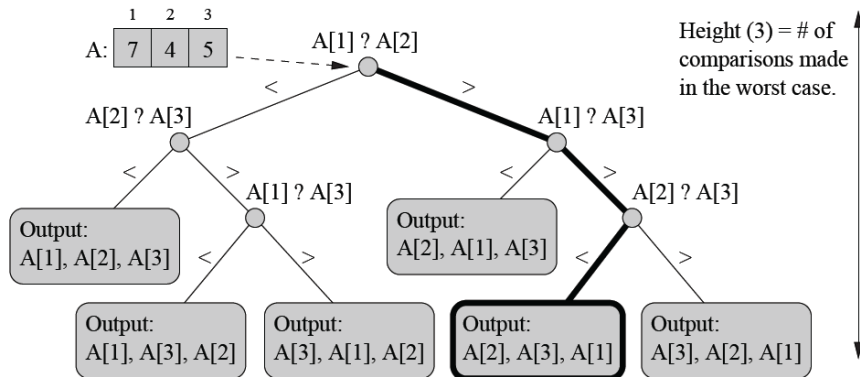
Lower Bounds for Comparison-Based Sorting Algorithms

- As it turns out, *any* algorithm that correctly sorts in the comparison-based model must make $\Omega(n \log n)$ comparisons in the worst case, so there is an $\Omega(n \log n)$ lower bound on the worst-case running time of any such algorithm.
- So an $O(n \log n)$ sorting algorithm has an **optimal** worst-case runtime in the comparison model.
- Much harder in general to prove lower bounds than upper bounds.
 - Most lower bounds we know are trivial: e.g., it take $\Omega(n^2)$ time to multiply two $n \times n$ matrices.

8

Decision Trees

- Any deterministic comparison-based algorithm can be modeled as an abstract *decision tree*:



9

Lower Bounds for Comparison-Based Sorting Algorithms

- A decision tree of height h has $\leq 2^h$ leaves
(assume for simplicity that we are sorting distinct elements, so $<$ and $>$ are the only two possible comparisons)
- Suppose # of leaves is less than $n!$...
 - Then the **pigeonhole principle** tells us that two different input permutations end up at the same leaf.
 - These inputs “look the same” to our algorithm, since they generate the same comparison results.
 - Our algorithm would apply the same permutation to both inputs, sorting one of them incorrectly!
- Therefore, we must have $2^h \geq n!$ for our algorithm to sort correctly. By Stirling’s approximation, we have $h \geq \log n! = \Theta(n \log n)$, so $h = \Omega(n \log n)$.

10

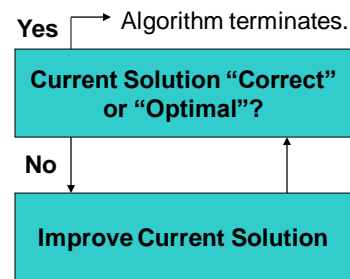

Lower Bounds Based on Reductions

- A similar decision tree argument can be used to prove that an algorithm for element uniqueness requires $\Omega(n \log n)$ worst-case time in the comparison model.
- Both sorting and element uniqueness have worst-case lower bounds of $\Omega(n \log n)$ in the real RAM model, but these can be harder to prove!
- Now that we have a non-trivial lower bound for these problems, we can carry this bound over to other problems via **reductions**.
 - Example: It takes $\Omega(n \log n)$ worst-case time to determine the closest pair out of n points in the plane.
(careful: this not a comparison-based problem!)

11

Bubble Sort (Iterative Refinement)


- Repeatedly scan array, swapping out-of-order pairs of consecutive elements.
- $O(n^2)$ time ($\Theta(n^2)$ in the worst case).
- Stable, in-place



12

Insertion Sort (Incremental Construction)

- In each iteration, enlarge sorted prefix of array by one element.
- Runtime $\Theta(n + I) = O(n^2)$, where I is the number of **inversions** in our input.
 - Good for sorting nearly-sorted sequences. In fact, if $I = O(n)$, then insertion sort runs in linear time!
- Stable, in-place.




Selection Sort

- Scan $A[1 \dots n]$ for smallest element, and swap it with $A[1]$.
- Then scan $A[2 \dots n]$ for smallest element, and swap it with $A[2]$. Etc.
- What is its running time? Is it stable and/or in-place?
- How would you classify this algorithm?
- How does it compare to bubble sort and insertion sort?

14

Merge Sort (Divide and Conquer)


- Recursively sort first and second halves of array, then **merge** the two resulting sorted subarrays in linear time.
(i.e., “divide, conquer, recombine”)
- Runtime $\Theta(n \log n)$, so it’s an optimal comparison-based sort.
- Stable, but not in-place.



15

Merge Sort Analysis: Running Time per Element of Data


- It takes $\Theta(n)$ time to merge two lists of combined length n .
- Think of this as $O(1)$ time per element taking part in the merge.
- Now look at a particular array element e . The total “work” we spend on e is equal to the number of merges in which e takes part: $O(\log n)$.
 - Why $O(\log n)$? Each merge doubles the size of the sorted subarray containing e .
- So $O(n \log n)$ total work.



16

Quicksort (Divide and Conquer)

- In linear time, **partition** array based on the value of some **pivot** element.
- Then recursively sort left side (all elements \leq pivot) and right side (\geq pivot).
- Typical implementation is in-place, but not stable.
- How to partition:
 - in linear time?
 - if some elements equal?



17

Quicksort Variants

- **Simple quicksort.** Choose pivot using a simple deterministic rule; e.g., first element, last element, median($A[1]$, $A[n]$, $A[n/2]$).
 - $\Theta(n \log n)$ time if “lucky”, but $\Theta(n^2)$ worst-case.
- **Deterministic quicksort.** Pivot on median (we’ll see shortly how to find the median in linear time).
 - $\Theta(n \log n)$ time, but not the best in practice.
- **Randomized quicksort.** Choose pivot uniformly at random.
 - $\Theta(n \log n)$ time with high probability, and fast in practice (competitive with merge sort).

18

Further Thoughts

- Any sorting algorithm can be made **stable** at the expense of **in-place** operation
(so we can implement quicksort to be stable but not in-place, or in-place but not stable).
- Memory issues:
 - Rather than sort large records, sort pointers to records.
 - Some advanced sorting algorithms only move elements of data $O(n)$ total times.
 - How will caching affect the performance of our various sorting algorithms?

19

Back to Comparison-Based Sorting – An Obvious Question...

Algorithm	Runtime	Stable	In-Place?
Bubble Sort	$O(n^2)$	Yes	Yes
Insertion Sort	$O(n^2)$	Yes	Yes
Merge Sort	$\Theta(n \log n)$	Yes	No
Randomized Quicksort	$\Theta(n \log n)$ w/high prob.	No*	Yes*
Deterministic Quicksort	$\Theta(n \log n)$	No*	Yes*
Heap Sort	$\Theta(n \log n)$	No*	Yes*

- Is stable in-place sorting possible in $O(n \log n)$ time in the comparison-based model?

* = can be transformed into a stable, out-of-place algorithm

20

The Ideal Comparison-Based Sorting Algorithm...

- ...would be **stable** and **in-place**.
- ...would require only **$O(n)$ moves** (memory writes)
- ...would be **simple** and **deterministic**.
- ...would run in **$O(n \log n)$** time.
- We currently only know how to achieve limited combinations of the above properties...

21

Adaptive Sorting

- The ideal sorting algorithm would run in $O(n \log n)$ worst-case time, but would be even faster if the input is “nearly sorted”
 - E.g., Merge sort, with its $\Theta(n \log n)$ running time, is much worse than insertion sort, in the event that we are sorting an already sorted array (which is not all that uncommon in practice...)
- How can we characterize unsortedness?

22

Adaptive Sorting




Figure 3. The partial order of the 11 measures of disorder. The ordering with respect to \leq_{alg} is up the page; for example, *Inv* optimality implies *Dis* optimality.

Table 1. The Worst-Case Lower Bounds for Different Measures of Disorder

M	Lower bound: $\log \text{below}(M(X), X , M) $
<i>Dis</i>	$\Omega(X (1 + \log [Dis(X) + 1]))$
<i>Exc</i>	$\Omega(X + Exc(X) \log [Exc(X) + 1])$
<i>Enc</i>	$\Omega(X (1 + \log [Enc(X) + 1]))$
<i>Inv</i>	$\Omega(X (1 + \log [\frac{Inv(X)}{ X } + 1]))$
<i>Max</i>	$\Omega(X (1 + \log [Max(X) + 1]))$
<i>Osc</i>	$\Omega(X (1 + \log [\frac{Osc(X)}{ X } + 1]))$
<i>Reg</i>	$\Omega(X + \log [Reg(X) + 1])$
<i>Rem</i>	$\Omega(X + Rem(X) \log [Rem(X) + 1])$
<i>Runs</i>	$\Omega(X (1 + \log [Runs(X) + 1]))$
<i>SMS</i>	$\Omega(X (1 + \log [SMS(X) + 1]))$
<i>SUS</i>	$\Omega(X (1 + \log [SUS(X) + 1]))$

From "A survey of adaptive sorting algorithms", by Estivill-Castro and Wood (1992)

23


RAM Sorting Algorithms

- Suppose we are sorting n integers in the range $0 \dots C - 1$ in the RAM model of computation.
- **Counting sort:** $O(n + C)$ time.
 - Sorts integers of magnitude $C = O(n)$ in linear time.
- **Radix sort:** $O(n \max(1, \log_n C))$ time.
 - Sorts integers of magnitude $C = O(n^k)$, $k = O(1)$, in linear time.

24

Counting Sort


- Scan $A[1..n]$ in $O(n)$ time and build an array $N[0..C-1]$ of element counts.
- By scanning N , we then reconstruct A in sorted order in $O(n + C)$ time.
- Ideally suited for $C = O(n)$.
- Not in-place.
- Stable, if we're careful...



25

Radix Sort

- Write elements of $A[1..n]$ in some base (radix), r . Typically, we set $r = n$.
- Sort on each digit, starting with the **least** significant, using a **stable** sort.
- # digits = $\log_n C$, which is constant if $C = n^{O(1)}$.
- Runtime $O(n)$ if $C = n^{O(1)}$. (recall word size assumptions w/RAM)
- Stable, not in-place



26

Implications of Sorting Lower Bounds: Might This Be Possible?

- Can we build a comparison-based priority queue satisfying these bounds?
 - Insert: $O(1)$ amortized time
 - Remove-Min: $O(1)$ amortized time

Implications of Sorting Lower Bounds: Might This Be Possible?

- Can we build a comparison-based priority queue satisfying these bounds?
 - Insert: $O(1)$ amortized time
 - Remove-Min: $O(1)$ amortized time
- Can we achieve these bounds with a comparison-based min-aware data structure (a data structure that always knows its minimum element):
 - Insert: $O(1)$ amortized time
 - Delete: $O(1)$ amortized time

Implications of Sorting Lower Bounds: Might This Be Possible?

- Can we achieve these bounds with a comparison-based min-aware data structure (a data structure that always knows its minimum element):
 - Insert: $O(1)$ time
 - Delete: $O(1)$ time*
- This doesn't seem to contradict the $\Omega(n \log n)$ lower bound on sorting, does it?
- These guarantees would be nice to have if we only occasionally need the functionality of a priority queue.

* Unless deleting the minimum, which takes $O(\log n)$ time.

Rank-Sensitive Priority Queues

- It turns out this is the “right” answer:
 - Insert: $O(\log(n/r))$ time
 - Delete: $O(\log(n/r))$ time
(r = rank of the element being inserted or deleted, where $r = 1$ is the minimum and $r = n$ is the maximum)
- So running time scales gracefully from $O(\log n)$ for operating on elements close to the minimum down to $O(1)$ for elements far away from the minimum (say, in the larger 99% of the data set).