


Lecture 8. Binary Search Trees

CpSc 8400: Algorithms and Data Structures
Brian C. Dean




School of Computing
Clemson University
Spring, 2016

Dictionary / Set Data Structures

- A dictionary maintains a set of elements with associated keys, supporting the operations:
 - *find*(*k*) : return a pointer to the element having key *k*, or indicates that such an element does not exist.
 - *insert*(*e*, *k*) : insert a new element *e* with key *k*.
 - *delete*(*e*) : delete element *e*, given a pointer to *e*.
To delete an element given its key: *delete*(*find*(*k*)).
- Common dictionary implementations:
 - Arrays and linked lists: not very efficient (at least one of the above operations takes $\Theta(n)$ worst-case time).
 - Balanced BSTs: all ops take $O(\log n)$ time.
 - Universal Hash tables: $O(1)$ (expected) time per op.

Binary Search Trees


- A BST is a binary tree satisfying the **binary search tree property**:



- Each node typically maintains a pointer to its parent, left child, and right child.
- Tree is **balanced** if height = $O(\log n)$.

3

Example




4

Fundamental Operations

Most BST operations can be implemented in a simple recursive fashion in $O(\text{height})$ time ($O(\log n)$ if balanced).
For example:

- **Find(T, k):**

- if $T == \text{NULL}$, then return NULL .
 - if $k == T.\text{key}$, then return T .
 - if $k < T.\text{key}$, then return $\text{Find}(T.\text{left}, k)$.
 - else return $\text{Find}(T.\text{right}, k)$.



- **Insert(T, e, k):**

- if $T == \text{NULL}$, then return $\text{NewNode}(e, k)$.
 - if $k < T.\text{key}$, then $T.\text{left} = \text{Insert}(T.\text{left}, e, k)$.
 - else $T.\text{right} = \text{Insert}(T.\text{right}, e, k)$.
- return T .


5

Tree Traversals

- We often enumerate the contents of an n -element BST in $O(n)$ time using an **inorder** tree traversal:

- Inorder(T):**

- if $T == \text{NULL}$, then return.
 - Inorder($T.\text{left}$).
 - print $T.\text{key}$.
 - Inorder($T.\text{right}$).



- Other types of common traversals:

- Preorder: root, left subtree, right subtree.
 - Postorder: left subtree, right subtree, root.
 - Eulerian: root, left subtree, root, right subtree, root.
(sequence of nodes encountered when we “walk around” a BST)

6


Sorting with a BST

- An inorder traversal prints the contents of an n -element BST **in sorted order** in $O(n)$ time.
- Therefore, we can use BSTs to sort: *insert* n elements, then do an inorder traversal.
- Since the BST is a comparison-based data structure, this means *insert* must run in $\Omega(\log n)$ time in the worst case; otherwise we could sort faster than $O(n \log n)$ in the comparison model.

7

Minimum and Maximum

- It's easy to locate the minimum and maximum elements in a BST in $O(h)$ time.
- For the minimum, start at the root and walk left as long as possible:




- Vice-versa for the maximum...


8

Predecessor and Successor

- The $pred(e)$ operation takes a pointer to e and locates the element immediately preceding e in an inorder traversal.



If e has a left child, then $pred(e)$ is the maximum element in e 's left subtree.



If e has no left child, then $pred(e)$ is the first "left parent" we encounter when walking from e up to the root.

- The successor operation, $succ(e)$, is completely symmetric to this.

9

Deletion

- It's easy to *delete* an element with zero or one children.
- To *delete* an element e with two children, first swap e with either $pred(e)$ or $succ(e)$, then proceed to delete e .
 - Note that if e has two children, then $pred(e)$ and $succ(e)$ can both have at most one child.
 - Also note that replacing e with $pred(e)$ or $succ(e)$ is o.k. according to the BST property.
- Can also replace e with the *join* of e 's two subtrees.

10

Using *pred* and *succ* for Inexact Searches

- In addition to *pred*(e) and *succ*(e) that take pointers to elements, we could implement:
 - *pred*(k) : find the element with largest key $\leq k$.
 - *succ*(k) : find the element with smallest key $\geq k$.
- This allows us to perform inexact searches: if an element with key k isn't present, we can at least find the closest nearby elements in the BST.
- This is one reason a BST is so powerful. Other dictionary data structures, notably hash tables, don't support *pred* / *succ* and cannot perform inexact searches.

11

Paging Through Nearby Elements

- Starting from element e, we can use *pred* / *succ* to find elements near e very quickly.
- Prototypical application: library catalog.
 - After searching for an author's name, you are presented with a alphabetized list in which you can scroll through nearby names.
- Can also answer **range queries**: output all elements whose keys are in the range [a, b].
 - Start at the element $e = \text{succ}(a)$. Then repeatedly call *succ*(e) until you reach b.
 - Running time $O(k + \log n)$, where $k = \#$ of elements written as output.

12

Rank-Based Access

- A BST augmented with subtree sizes can support two additional useful operations:
 - *select(k)* : return a pointer to the *k*th largest element in the tree (*k* = 1 is the min, *k* = *n* is the max, and *k* = *n*/2 is the median).
 - *rank(e)* : given a pointer to *e*, returns the number of elements with keys $\leq e$'s key.
(that is, *e*'s index within an inorder traversal).
- The element of rank *k* in a set is also called the ***k*th order statistic** of the set. Accordingly, the CLRS textbook calls a BST augmented this way an **order statistic tree**.

13

Select and Rank

Select(*T*, *k*):

$r = \text{size}(T.\text{left}) + 1.$ ($r = \text{rank of root}$)

if $k = r$, then return *T*.


if $k < r$, then return *Select*(*T*.left, *k*).

else, return *Select*(*T*.right, $k - r$).

Rank(*e*):

Add up the total # of yellow elements as we walk from *e* up to the root.

(these are all the elements less than *e*)



14


Dynamic Sequences

- Suppose we use an array or linked list to encode a sequence of n elements.
 - We can insert/delete at the endpoints of the sequence in $O(1)$ time,
 - But insertions/deletions in the middle of the sequence take $O(n)$ worst-case time.
- Using a balanced BST augmented with subtree sizes, we can perform all the following operations in $O(\log n)$ time:
 - Insert / delete anywhere in the sequence.
 - Access or modify any element by its index in the sequence.

15

Dynamic Sequences

- The “BST property” is slightly different when we use a BST to encode a sequence $A[1..n]$:




- Elements no longer have keys, and we no longer use *find*. Rather, we rely on the *select* to access elements based their index within the sequence.

16

Dynamic Sequences : Example

Sequence: 3141592653589



17


Dynamic Sequences

- Many operations have a nice meaning:
 - *select(i)* accesses the *i*th element
 - An inorder traversal prints the sequence in order
 - *succ* and *pred* allow us to move between successive elements (as in a linked list)
 - *min* and *max* jump to the beginning and end
- You might want to think of a BST as a structure that fundamentally encodes an arbitrary sequence from left to right. In the case of a dictionary, this sequence happens to be the sorted ordering of the elements we are storing.

18

Balancing a BST

- There are many ways to modify the insert and delete operations of a BST to keep it balanced:
 - **Worst-case** mechanisms: $h = O(\log n)$ always.
AVL trees, BB[α] trees, red-black trees
 - **Amortized** mechanisms: starting with an empty tree, any sequence of m operations takes $O(m \log n)$ time.
splay trees, “batch” rebalancing methods
 - **Randomized** mechanisms: always balanced “with high probability”. *randomly-balanced BSTs, treaps*
- Most of these are based on **rotations**:



19

Example of Balancing with Rotations: AVL Trees

- AVL: Adel'son-Vel'skiĭ, and Landis '62.
- A binary tree is **height balanced** if:
 - Its left and right subtrees differ in height by at most 1, and
 - Its left and right subtrees are themselves height balanced.
- It's easy to show that height-balanced trees are balanced (proof on next slide).
- An AVL tree is a height-balanced tree where every node is augmented with the height of its subtree.
 - Note: easy to keep this info up to date after a rotation.
- After each insertion or deletion, we restore the height balance property by performing $O(\log n)$ carefully chosen rotations.

20


Height Balanced \rightarrow Balanced

- **Claim:** A height-balanced tree of height h contains $\geq F_h$ nodes, where F_h is the h^{th} Fibonacci number.
- Recall that $F_h = \Theta(\Phi^h)$, where $\Phi = (1+\sqrt{5})/2$. Therefore, an n -element height-balanced tree can have height at most $O(\log n)$.
- Easy proof by induction:
 - Consider any arbitrary height-balanced tree of height h .
 - Without loss of generality, suppose its left subtree is tallest.
 - By induction, left subtree contains $\geq F_{h-1}$ elements.
 - By induction, right subtree contains $\geq F_{h-2}$ elements.
 - So total # of elements is $\geq 1 + F_{h-1} + F_{h-2} = 1 + F_h$.
 - Don't forget the base cases:
 - Tree of height 0 contains ≥ 1 element ($F_0 = 0$).
 - Tree of height 1 contains ≥ 2 element ($F_1 = 1$).

21

Restoring Balance After an Insertion

- After an *insert*, walk back up the insertion path and increment augmented subtree heights.
- Stop at the first node (z in the figure below) at which we encounter a height-balancing violation (a height imbalance of exactly 2).




- Using 1 or 2 rotations, we will be able to rebalance our tree at z .
- The height of z 's subtree will decrease by 1 to its original value, so we don't need to visit the remaining nodes on the insertion path above z .

22

Restoring Balance After an Insertion

- Let's suppose we have an imbalance at node z due to an insertion that makes z 's left subtree too tall (the same argument will apply to the right subtree as well...)
- Easy case:** Insertion into subtree A makes it too tall:




- Harder case:** Insertion into subtree B makes it too tall...

23

Restoring Balance After an Insertion

- Harder case:** Subtree B too tall. There are now 2 possible subcases:



- In all both cases, a left rotation about the edge (x, y) brings us back to the previous case (with A being too tall).
- So 2 total rotations needed for these cases.

24