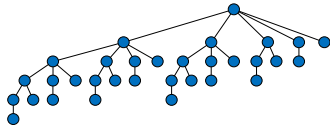


Lecture 11. Randomized Algorithms

CpSc 8400: Algorithms and Data Structures
Brian C. Dean



School of Computing
Clemson University
Spring, 2016

Randomized Algorithms

- Randomized algorithms offer many advantages:
 - Simplicity (e.g., randomly-balanced BSTs).
 - Malicious adversary often cannot force worst-case performance.
 - Efficiency (e.g., randomized quicksort vs. deterministic $O(n \log n)$ quicksort)
 - In some cases, randomization even gives us more power than we have with deterministic algorithms!
- ... and some disadvantages:
 - May be more difficult to debug.
 - Analysis requires some background in probability.
 - We need a good source of randomness...


Randomized Algorithms

- Two types of randomized algorithms:
 - **Las Vegas**. Always correct but running time may vary depending random choices made by algorithm.
 - **Monte Carlo**. Deterministic running time but may make mistakes.
- We can convert any Las Vegas algorithm into a Monte Carlo algorithm by stopping it early if it fails to run fast enough.
 - So Las Vegas often viewed as the “better” of the 2 types.
- Can convert Monte Carlo to Las Vegas if we can check correctness of a solution: keep running until we get a correct solution.

3

Analyzing Randomized Algorithms

- With a Monte Carlo algorithm, we typically want to prove that the output is correct with high probability.
- With Las Vegas algorithms, we typically want to show that a certain running time bound holds with high probability.
 - Another popular analysis goal (which we’ll talk about soon), is to compute the “expected” running time of the algorithm.



4

Average-Case Analysis Versus Randomized Algorithms

- Recall that we usually focus on **worst-case** performance of algorithms.
- Sometimes (especially when worst-case inputs do not often occur in practice), it makes sense also to consider **average-case** performance – the expected running time when given a random input (according to some probability distribution that mimics what you expect in practice).
- However, note that this is **completely different** from analyzing the (worst-case) expected running time of a randomized algorithm.
 - Randomness in input vs. randomness in algorithm.

5

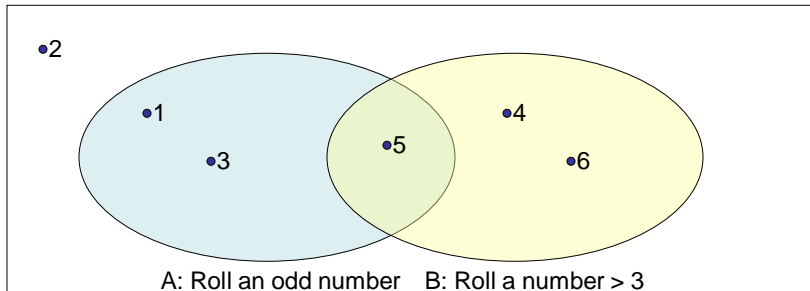
Basic Probability Theory

- A random trial is described by a set of **outcomes**, each with an associated **probability**.
- Probabilities of all outcomes sum to 1.
- An **event** is a set of outcomes, and the probability of an event is the sum of the probabilities of these outcomes.
- Example: Let E be the event that we roll two dice and they sum to 4.
 - This includes the outcomes (1, 3), (2, 2), and (3, 1) each with probability 1/36, so $\Pr[E] = 3/36 = 1/12$.
- $\Pr[\sim E] = 1 - \Pr[E]$.

6

Intersections and Unions of Events

- If A and B are events, then
 - $A \cap B$ is the event that both occur
 - $A \cup B$ is the event that either A or B occur, or both.
- $\Pr[A \cup B] = \Pr[A] + \Pr[B] - \Pr[A \cap B]$, as we can see from **Venn diagram** below (for rolling a 6-sided die).



7

Back to Randomized Quicksort...

- If we select pivot elements that result in “balanced partitions” (e.g., the median), quicksort runs in $\Theta(n \log n)$ time.
- If we select pivot elements poorly, quicksort could run in $\Theta(n^2)$ time.
- **Claim:** Randomized quicksort (where we select pivot elements at random) runs in $O(n \log n)$ time with high probability.
 - But how do we prove this...?
 - And what does “with high probability” mean?

8

“With High Probability”

- A Monte Carlo randomized algorithm produces correct output **with high probability** if

$$\Pr[\text{failure}] \leq 1 / n^c$$

where c can be chosen to be any constant we wish, by adjusting the hidden constant in the $O(\dots)$ running time of the algorithm.

- A Las Vegas randomized algorithm runs in $O(X)$ time **with high probability** if

$$\Pr[\text{fails to run in } O(X) \text{ time}] \leq 1 / n^c$$

where c can be chosen to be any constant we wish, by adjusting the hidden constant in the $O(X)$ running time.

9

Intersection of Events and Independence

- $\Pr[A \cap B] = \Pr[A] \Pr[B]$ if and only if events A and B are **independent**.
- Two events are independent if knowledge of the occurrence of one event has no impact on the probability of the occurrences of the other event.
- Example: Boosting success probability via repeated trials.
 - Take a Monte Carlo algorithm that fails with probability $\leq 1/2$ (and suppose we can detect when it fails).
 - If we run it k times, the probability of failure drops to $\leq 1/2^k$.
 - If we run it $k = c \log n$ times, the probability of failure drops to $\leq 1/n^c$ (i.e., the algorithm succeeds with high probability)

10

Reducing Problem Size by a Constant Fraction per Iteration

- Suppose we have an algorithm for which:
 - We start with a problem of size n .
 - In every iteration, the effective size of the problem is reduced to a constant fraction of its original size.
- Then, it's easy to see that our algorithm must perform only $O(\log n)$ iterations.
- Prototypical example: binary search.
 - Each iteration reduces problem to $\frac{1}{2}$ of its original size.

11

The Randomized Reduction Lemma

- Suppose we have an algorithm for which:
 - We start with a problem of size n .
 - In every iteration, the effective size of the problem is reduced to a constant fraction of its original size **with some constant probability**.
- Then, our algorithm performs only $O(\log n)$ iterations **with high probability**.
- We call this the **randomized reduction lemma**, and it will be one of the main tools we use for analyzing randomized algorithms and data structures.

12


Example: Randomized Binary Search

- Suppose we're searching a sorted array $A[1 \dots n]$ for an element of value X .
- In a standard binary search, we compare $A[n/2]$ and X , then recursively search either $A[1 \dots n/2]$ or $A[n/2 \dots n]$.
- Suppose instead of $A[n/2]$, we compare X against a random element and then recursively search left or right as before...

13

Example: Randomized Binary Search

- **Claim:** Randomized binary search runs in $O(\log n)$ time with high probability.
- **Simple proof:** If we happen to choose a "pivot" element $A[j]$ where $n/3 \leq j \leq 2n/3$ (and this happens with probability $\geq 1/3$), then our problem will be reduced to $\leq 2/3$ of its original size.



14

The Union Bound

- Recall that $\Pr[A \cup B] = \Pr[A] + \Pr[B] - \Pr[A \cap B]$.
- It typically suffices just to use the rough upper bound $\Pr[A \cup B] \leq \Pr[A] + \Pr[B]$.
- For multiple events $E_1 \dots E_k$, this gives us what is known as the **union bound** or **Boole's inequality**:
 $\Pr[E_1 \cup E_2 \cup \dots \cup E_k] \leq \Pr[E_1] + \dots + \Pr[E_k]$.
- For example:
 - Suppose each of 50 parts in a complex machine fails with probability $\leq 1/100$
 - Then $\Pr[\text{entire machine fails}] \leq 50(1/100) = 1/2$.

15

The Union Bound

- If there are n bad events that can happen, and each happens with probability $\leq p$, then the probability **any** bad event happens is at most np .
I.e, for events $E_1 \dots E_n$, $\Pr[\cup_i E_i] \leq \sum_i \Pr[E_i]$
- This meshes particularly well with our definition of “with high probability”: If some property holds for a generic input element w.h.p., then it also holds for each of our n input elements w.h.p.
- **Example:** If randomized quicksort spends only $O(\log n)$ on a generic input element w.h.p, then its total running time is $O(n \log n)$ w.h.p.

16

A Prototypical “High Probability” Analysis...

- **Step 1:** Consider some generic input element e . Show that randomized quicksort spends $O(\log n)$ work on e w.h.p.
 - Usually easy with the randomized reduction lemma.
- **Step 2:** The union bound automatically allows us to extend this result to show that randomized quicksort spends $O(\log n)$ work on **each** of its input elements w.h.p.
- So total running time is $O(n \log n)$ w.h.p.

17

Randomized Reduction and Randomized Quicksort

Randomized binary search for 7

1 2 3 7 8 9 12 13 14 15

1 2 3 7 8 9

3 7 8 9

3 7

 = random choices of “pivots”

Randomized quicksort

3 7 9 2 8 1 12 14 15 13

1 2 8 7 9 3 13 14 15

1 7 3 8 9 13 15

3 7 9


3

 = random choices of pivots

18

Quicksort and BST Construction

- There is a direct analog between quicksort and the process of building a BST!
- Hence, both have the same running times.
- Randomized quicksort is analogous to building a BST by inserting elements in random order.



19

Randomly-Built BSTs

- **Theorem:** if we build a BST on n elements by inserting them in random order, then with high probability each call to insert will take $O(\log n)$ time.
- Equivalently, with high probability:
 - Each element will have depth $O(\log n)$.
 - The entire tree will have depth $O(\log n)$.
 - The entire tree will take $O(n \log n)$ time to build.
- **Corollary:** randomized quicksort runs in $O(n \log n)$ time with high probability!

20

Randomly-Built BSTs

- **Theorem:** if we build a BST on n elements by inserting them in random order, then with high probability each call to insert will take $O(\log n)$ time.
- Equivalently, with high probability:
 - Each element will have depth $O(\log n)$.
 - The entire tree will have depth $O(\log n)$.
 - The entire tree will take $O(n \log n)$ time to build.
- **Corollary:** randomized quicksort runs in $O(n \log n)$ time with high probability!

21

Randomized Quicksort, Binary Search, and BST Construction

Randomized binary search for 7

1 2 3 7 8 9 12 13 14 15

1 2 3 7 8 9

3 7 8 9

3 7

 = random choices of "pivots"

Randomized quicksort / BST Construction

3 7 9 2 8 1 12 14 15 13

1 2 8 7 9 3 13 14 15

1 7 3 8 9 13 15

3 7 9

3

 = random choices of pivots

22

Maintaining Randomness


- If we build a BST at random on n elements, then with high probability it will be balanced.
- However, subsequent calls to insert and delete might cause the tree to become unbalanced.
- Remarkably, we can fix this by doing some carefully chosen random rotations after each insert and delete so the tree is always in a state that is “as if it was just randomly built from scratch”.
 - More precisely: within each subtree, each element is equally likely to be at the root.
- This gives us a simple randomized mechanism for keeping a BST balanced with high probability.

23


Randomly-Balanced BSTs

- To insert an element e into an $(n - 1)$ -element tree:
 - With probability $1/n$, insert e at the root (insert as usual, then rotate up to root).
 - Otherwise (with probability $1 - 1/n$), recursively insert into the left or right subtree of the root,
- To delete an element e , replace e with the a **randomized join** of e 's two subtrees L and R :

With probability $|L| / (|L| + |R|)$:



With probability $|R| / (|L| + |R|)$:



24