


## Lecture 12. Randomized Data Structures

**CpSc 8400: Algorithms and Data Structures**  
**Brian C. Dean**



**School of Computing**  
**Clemson University**  
**Spring, 2016**

### Warm-Up: Randomized Reduction Practice

- Take an array of distinct numbers and randomly permute it.
- If you scan through the array keeping a running maximum, how many times will this maximum be reset?

## Recap: The Randomized Reduction Lemma

- Suppose we have an algorithm for which:
  - We start with a problem of size  $n$ .
  - In every iteration, the effective size of the problem is reduced to a constant fraction of its original size **with some constant probability**.
- Then, our algorithm performs only  $O(\log n)$  iterations **with high probability**.

3


## Warm-Up: Randomized Reduction Practice

- Take an array of distinct numbers and randomly permute it.
- If you scan through the array keeping a running maximum, how many times will this maximum be reset?
- How many non-dominated points do we expect to see in a collection of  $n$  random points in 2D?

4

## Warm-Up: Randomized Reduction Practice


- Take an array of distinct numbers and randomly permute it.
- If you scan through the array keeping a running maximum, how many times will this maximum be reset?
- How many non-dominated points do we expect to see in a collection of  $n$  random points in 2D?



5

## Simple Application: Mergeable Heaps


- Suppose we store elements in heap-ordered binary trees.
- All priority queue operations easy once we have *merge*:
  - Insert: merge with a new 1-element heap
  - Remove-min: remove root, merge two child subtrees back together
- A simple way to merge is using randomization!
  - Take two heap-ordered trees  $h_1$  and  $h_2$ ,  $h_1$  having the smaller root.
  - Clearly,  $h_1$ 's root must become the root of the merged tree.
  - To complete the merge, recursively merge  $h_2$  into either the left or right subtree of  $h_1$ :



6

## Simple Application: Mergeable Heaps


- Suppose we store elements in heap-ordered binary trees.
- All priority queue operations easy once we have *merge*:
  - Insert: merge with a new 1-element heap
  - Remove-min: remove root, merge two child subtrees back together
- A simple way to merge is using randomization!
  - Take two heap-ordered trees  $h_1$  and  $h_2$ ,  $h_1$  having the smaller root.
  - Clearly,  $h_1$ 's root must become the root of the merged tree.
  - To complete the merge, recursively merge  $h_2$  into either the left or right subtree of  $h_1$ :



7


## Merging Two Heap-Ordered Trees (Null Path Merging Viewpoint)

- A **null path** is a path from the root of a tree down to an “empty space” at the bottom of the tree.
- Given specific null paths in  $h_1$  and  $h_2$ , it's easy to merge  $h_1$  and  $h_2$  along these paths.
  - The keys along a null path are a sorted sequence.
  - Merging along null paths is like merging two sorted sequences.
  - This process is also equivalent to the recursive merging process from the previous slide.



8

## Merging Two Heap-Ordered Trees (Null Path Merging Viewpoint)



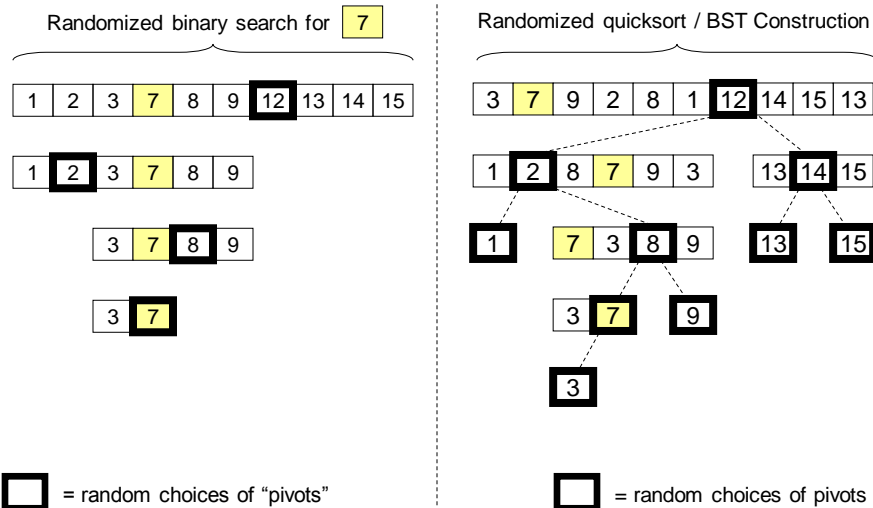
9

## Running Time Analysis

- The time required to merge two heaps along null paths is proportional to the combined lengths of these paths.
- So all we need is a method to find “short” null paths and we will have an efficient merging algorithm.
- Note that every  $n$ -node binary tree has a null path of length  $O(\log n)$ .
- There are many ways to find short null paths, each of which leads us to a different mergeable heap data structure (e.g., leftist heaps, skew heaps).
- Our preceding approach is perhaps the simplest – just choose null paths by “walking down the tree randomly”!

10

## Recap: Randomized Quicksort, Binary Search, and BST Construction



11

## Recap: Randomly-Built BSTs


- **Theorem:** if we build a BST on  $n$  elements by inserting them in random order, then with high probability each call to insert will take  $O(\log n)$  time.
- Equivalently, with high probability:
  - Each element will have depth  $O(\log n)$ .
  - The entire tree will have depth  $O(\log n)$ .
  - The entire tree will take  $O(n \log n)$  time to build.
- **Corollary:** randomized quicksort runs in  $O(n \log n)$  time with high probability!

12


## Recap: Randomly-Balanced BSTs

- To insert an element  $e$  into an  $(n - 1)$ -element tree:
  - With probability  $1/n$ , insert  $e$  at the root (insert as usual, then rotate up to root).
  - Otherwise (with probability  $1 - 1/n$ ), recursively insert into the left or right subtree of the root,
- To delete an element  $e$ , replace  $e$  with the a **randomized join** of  $e$ 's two subtrees  $L$  and  $R$ :

With probability  $|L| / (|L| + |R|)$ :




With probability  $|R| / (|L| + |R|)$ :



13

## Treaps



- A treap is a binary tree in which each node contains two keys, a “**heap key**” and a “**BST key**”.
- It satisfies the heap property with respect to the heap keys, and the BST property with respect to the BST keys.
- The BST keys will store the elements of our BST; we’ll choose heap keys as necessary to assist in balancing.

14

## Treaps

- If heap keys are all distinct, then there is only one valid “shape” for a treap. (why?)
- If we choose heap keys at random, then our treap will be randomly-structured!

What about *insert* and *delete*?

- *insert* : Insert new element as leaf using the standard BST insertion procedure (so BST property satisfied). Assign new element a random heap key. Then restore heap property (while preserving BST property) using *sift-up* implemented with rotations.
- *delete* : Give element a heap key of  $+\infty$ , sift it down (again using rotations, to preserve BST property) to a leaf, then delete it. Or, use *join*...

15

## Several Data Structures Satisfy a Combination of BST + Heap Properties

- **Treaps.** Elements stored in BST keys. Heap keys chosen randomly to help with balance.
- **Rank-Sensitive Priority Queues** [Dean, Jones '09]. Elements stored in heap keys. BST keys randomly to help with “balance”.
- **Priority Search Trees.** Holds a collection of points in the 2D plane; useful for 2D range searching. X coordinates form “BST” part, and y coordinates form “heap” part.
- **Cartesian Trees.** Represents a sequence in the BST part; the same elements also satisfy the heap property. Useful for a wide range of applications like range minimum queries, suffix tree construction, lowest common ancestors, and more.

16




## Skip Lists

- A skip list is a simple randomized dictionary data structure that provides  $O(\log n)$  w.h.p. performance guarantees (more or less equivalent to a randomly-balanced BST).
- Very simple to implement and analyze.
- Based on linked lists rather than BSTs (often billed as an alternative to balanced BSTs...)
- Suppose we store a dictionary as a sorted linked list. Recall that scanning down the list is the bottleneck operation (taking  $O(n)$  time in the worst case).
  - How might we speed this up?

17

## Example



- We insert a dummy “start” element (with effective key value  $-\infty$ ) that is present on all levels.
- Define  $L$  as the maximum level in the skip list.

18

## Fundamental Operations

- To *find* an element, repeatedly scan right until the next step would take us too far, then step down.
- To *insert* a new element,
  - Insert into the level-0 list.
  - Flip a fair coin. If heads, also insert into the level-1 list, then flip another fair coin, and if heads again, insert into level-2 list, etc.
- To *delete*, simply remove an element from every level on which it exists.
- Other operations like *pred*, *succ*, *min*, *max*, *rank*, and *select*, are easy to implement.

19


## Analysis

- The running time of each operation is dominated by the running time of finding an element, so let's focus on analyzing the running time of *find*.
- **Clever idea:** work backwards!
  - Starting from some element  $e$  in the level-0 list, retrace the *find* path in reverse.
  - Step up whenever possible, otherwise step to the left.
  - **Claim:** This runs in  $O(\log n)$  time w.h.p.
  - **Proof:** We'd like to use the randomized reduction lemma. But how...
  - Using the union bound, we can extend this to an  $O(\log n)$  w.h.p. for all elements in the skip list.  
(and this also shows that  $L = O(\log n)$  w.h.p.)

20

## Analysis

- Let  $N$  denote the number of elements in current level to the left of our current location during the backward scan.



- Claim:** In each step,  $N$  is reduced to half of its current value with probability at least  $\frac{1}{4}$ .

21

## Analysis

- Claim:** In each step,  $N$  is reduced to half of its current value with probability at least  $\frac{1}{4}$ .
- Proof:** For  $N$  to be reduced to  $\leq \frac{1}{2}N$  two events must occur:
  - A:** Our next step moves up, since we'd flipped heads at the current (element, level).
  - B:** At most half the  $N$  elements to our left in the current level also flipped heads (and hence also exist on the next level).
- $\Pr[A \cap B] = \Pr[A]\Pr[B]$  since  $A$  &  $B$  are independent.
- $\Pr[A] = \frac{1}{2}$  (unbiased coin flip)
- $\Pr[B] = \Pr[\text{at most } N/2 \text{ heads in } N \text{ coin flips}] \geq \frac{1}{2}$ .
- Hence,  $\Pr[N \text{ reduced to } \leq N/2] = \Pr[A \cap B] \geq \frac{1}{4}$ .

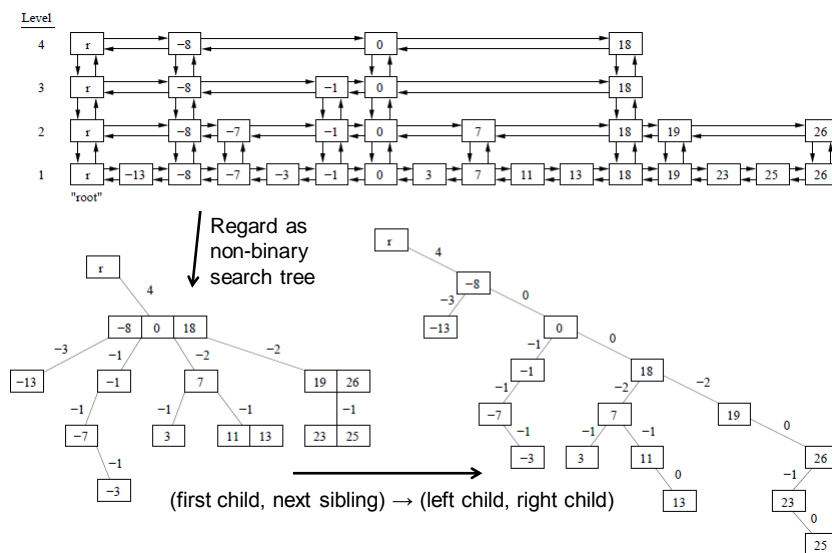
22

## Skip Lists Versus BSTs

- Skip lists were initially introduced as a “simpler alternative” to balanced BSTs.
  - Recall they provide more or less exactly the same functionality, but are based on linked lists, and do not involve any complicated tree-balancing mechanism.
- [Dean, Jones '07]: However, the skip list can also be converted into a fairly simple randomized BST balancing mechanism.
  - And one can even map the other way, converting any “BST” balancing mechanism to an equivalent “skip list” implementation. So if you want, say, static optimality in a skip list, use a “splay skip list”!

23

## Skip Lists Versus BSTs



24

## Operations

- Rebalancing after insert or delete is accomplished by a simple sequence of rotations, mimicking what happens with the skip list.

INSERT( $T, x$ ):

```
1 Insert  $x$  as a leaf in  $T$  (standard BST insert)
2 Set  $w(x, \text{parent}(x)) = H - h(\text{parent}(x))$ 
3 while Random(0,1) = 0
4   Promote( $x$ )
5 If  $w(x, \text{parent}(x)) = 0$  and  $x$  is a left child
6   Rotate  $x$  with  $\text{parent}(x)$ 
```

PROMOTE( $x$ ):

```
1 while  $w(x, \text{parent}(x)) = 0$ 
2   Rotate  $x$  with  $\text{parent}(x)$ 
3 Increment  $w(x, \text{parent}(x))$ 
4 Decrement  $w(x, \text{lchild}(x))$  and  $w(x, \text{rchild}(x))$ 
```

DELETE( $x$ ):

```
1 If  $x$  has no children, simply delete  $x$ 
2 while  $x$  has two children
3   Demote( $x$ )
4 Let  $c$  be the single child of  $x$ 
5  $w(x, \text{parent}(x)) = w(x, \text{parent}(x)) + w(x, c)$ 
6 Replace  $x$  with  $c$ .
```

DEMOTE( $x$ ):

```
1 if  $w(x, \text{rchild}(x)) = 0$ 
2   Rotate  $x$  with  $\text{rchild}(x)$ 
3 Decrement  $w(x, \text{parent}(x))$ 
4 Increment  $w(x, \text{lchild}(x))$  and  $w(x, \text{rchild}(x))$ 
5 while  $w(x, \text{lchild}(x)) = 0$ 
6   Rotate  $x$  with  $\text{lchild}(x)$ 
```