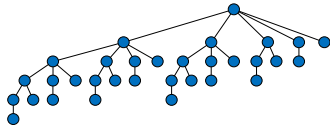


Lecture 14. Tail Bounds, Randomized Incremental Construction (Also Some Computational Geometry)

CpSc 8400: Algorithms and Data Structures
Brian C. Dean



School of Computing
Clemson University
Spring, 2016


Expected Value

- The expected value of a discrete random variable X , denoted $\mathbf{E}[X]$, is defined as

$$\mathbf{E}[X] = \sum_{\text{values } v} v \Pr[X = v].$$

- Think of $\mathbf{E}[X]$ informally as the “center of mass” of X ’s probability distribution.
- Example: Let D be the max of two dice rolls. Recall that D has this probability distribution.

– Thus, $\mathbf{E}[D] = 1(1/36) + 2(3/36) + 3(5/36) + 4(7/36) + 5(9/36) + 6(11/36) = 161/36 = 4 \frac{17}{36}$



- Careful: Don’t write $\mathbf{E}[A]$ if A is an event (another syntax error!)

Computing Expected Values

There are generally 4 different ways we will compute expected values in this class:

1. Directly using the definition $E[X] = \sum_v v \Pr[X = v]$.
2. The special case of an **indicator** random variable.
3. The special case of a **geometric** random variable.
4. Expressing a complicated random variable in terms of a sum of simpler r.v.'s and applying **linearity of expectation**.

3

Randomized Quicksort Revisited

- Let's think again about randomized quicksort.
- We've already shown an $O(n \log n)$ running time both w.h.p and also in expectation.
- There is another alternate expected running time proof that corresponds exactly to the w.h.p. proof:
- **W.h.p. proof:**
 - Randomized reduction lemma $\rightarrow O(\log n)$ / element w.h.p.
 - Union bound $\rightarrow O(n \log n)$ time for all elements w.h.p.
- **Expected running time proof:**
 - Linearity of expectation $\rightarrow O(\log n)$ expected time / element
 - Linearity of expectation $\rightarrow O(n \log n)$ total expected time.

4

Wald's Theorem

- Linearity of expectation tells us that for N identically distributed random variables $X_1 \dots X_N$,

$$E[X_1 + \dots + X_N] = N E[X_1].$$

- What if N itself is a random variable though? (i.e., we have a random number of trials, each of which have identically-distributed random running times)

- Wald's theorem** says that, as we might expect:


$$E[X_1 + \dots + X_N] = E[N] E[X_1].$$

(this requires a few technical conditions to hold; for example, N must be independent of the X_i 's)

5

Tail Bounds

- We are often interested in the probability that a random variable X deviates significantly from its mean, $E[X]$, when we randomly sample X from its distribution.
- Examples:
 - Show that $\Pr[T = O(\log n)] \geq 1 - 1/n^c$ for arbitrary $c > 0$.
 - What is $\Pr[X \geq 75]$ if X denotes the number of heads in 100 flips of a fair coin.



6

Markov's Inequality

- X is any *nonnegative* random variable, then
$$\Pr[X \geq k\mathbf{E}[X]] \leq 1/k.$$
- Sometimes written as $\Pr[X \geq a] \leq \mathbf{E}[X] / a$.
- Example: let X be the number of heads we see when flipping 100 coins.
 - $\mathbf{E}[X] = 50$.
 - $\Pr[X \geq 75] = \Pr[X \geq (3/2)\mathbf{E}[X]] \leq 2/3$.
- Due to its generality, Markov's inequality is a rather weak bound, although it's still quite useful.
- If expected running time = T , then probability our algorithm takes $\geq kT$ time is at most $1/k$.

7

Chernoff Bounds

- Suppose $X = X_1 + X_2 + \dots + X_n$ is a sum of independent indicator (0/1) random variables (a common scenario).
- Then the tails of X 's distribution drop off very quickly:
 - $\Pr[X \geq \mathbf{E}[X] + t] \leq e^{-2t^2/n}$
 - $\Pr[X \leq \mathbf{E}[X] - t] \leq e^{-2t^2/n}$
 - $\Pr[X \leq (1 - \epsilon)\mathbf{E}[X]] \leq e^{-\epsilon^2\mathbf{E}[X]/2}$
 - $\Pr[X \geq (1 + \epsilon)\mathbf{E}[X]] \leq e^{-\epsilon^2\mathbf{E}[X]/4}$ (if $\epsilon \leq 2e - 1$)
 - $\Pr[X \geq (1 + \epsilon)\mathbf{E}[X]] \leq 2^{-(\epsilon+1)\mathbf{E}[X]}$ (if $\epsilon > 2e - 1$)
- Example: If X denotes the number of heads we see when flipping 100 fair coins, then
 - $\Pr[X \geq 75] = \Pr[X \geq \mathbf{E}[X] + 25] \leq e^{-25^2/50} = e^{-12.5} \leq 0.000004$
- Much stronger than Markov's inequality, but only applies when X is a sum of independent indicator random variables.

8

Proving the Randomized Reduction Lemma

- **Theorem:** Suppose every iteration of our algorithm has probability $\geq p$ of reducing our problem size to at most $\leq q$ times its original size. Then given any constant $c > 0$, we can find another constant k such that $\Pr[T \geq k \log n] \leq 1 / n^c$, where T is the total number of iterations of our algorithm.
- **Proof:**
 - Suppose we run our algorithm for $k \log n$ iterations.
 - Let X denote the # of “good” iterations among these (a good iteration is one where we reduce the problem size to $\leq q$ times original).
 - Now note that $\Pr[T \geq k \log n] \leq \Pr[X \leq \log_{(1/q)} n]$, since the event that “ $T \geq k \log n$ ” is a subset of the event that “ $X \leq \log_{(1/q)} n$ ”.
 - So now let’s find a constant k such that $\Pr[X \leq \log_{(1/q)} n] \leq 1/n^c$.
 - Continued...

9

Proving the Randomized Reduction Lemma

- **Proof Continued:**
 - We run our algorithm for only $L = k \log_2 n$ iterations.
 - $X = \#$ of good iterations among these.
 - We want to show that $\Pr[X \leq \log_{(1/q)} n] \leq 1/n^c$.
 - Write $X = X_1 + X_2 + \dots + X_L$, where X_j is an indicator r.v. whose value is 1 if the j th iteration is good. Note that $\mathbf{E}[X_j] \geq p$.
 - $\mathbf{E}[X] = \mathbf{E}[X_1] + \dots + \mathbf{E}[X_L] \geq Lp = kp \log_2 n$.
 - So $\Pr[X \leq \log_{(1/q)} n]$
 $\leq \Pr[X \leq (1 / kp \log_2(1/q)) \mathbf{E}[X]]$
 $= \Pr[X \leq (r / k) \mathbf{E}[X]]$, where $r = 1 / [p \log_2(1/q)]$ is a constant.
 - Now we use the Chernoff bound with $(1 - \epsilon) = r / k$:
 $\Pr[X \leq (1 - \epsilon)\mathbf{E}[X]] \leq e^{-\epsilon^2 \mathbf{E}[X]/2} = e^{-(1-r/k)^2 \mathbf{E}[X]/2} \leq e^{-\mathbf{E}[X]/8}$, if we choose k sufficiently large so that $1 - r / k \geq 1/2$.
 - Thus, $\Pr[X \leq \log_{(1/q)} n] \leq e^{-\mathbf{E}[X]/8} \leq e^{-(kp/8) \log n} = e^{-(kp/8 \ln 2) \ln n} = 1 / n^{kp/8 \ln 2}$.
 - So $\Pr[X \leq \log_{(1/q)} n] \leq 1/n^c$ if we choose $k \geq (8 \ln 2)/p$.

10


Randomized Incremental Construction

- When building a BST or a binary heap, inserting elements in certain orderings may lead to inefficient running times.
- However, inserting in a random order is (with high probability) quite efficient.
- These are examples of a general algorithm design technique called **randomized incremental construction**, where we build a solution by incorporating one element at a time, processing elements in random order.
- R.I.C. algorithms are especially common in computational geometry:
 - Examples: Convex hulls (2D and 3D), half-space intersections, Voronoi diagrams, Delaunay triangulations, smallest enclosing circle, low-dimensional linear programming, binary space partition trees, trapezoidal decompositions, closest pair, etc.

11

Example: Smallest Enclosing Circle

- Given n points in the plane, find the smallest circle enclosing them all.



12

Smallest Enclosing Circle

- Given n points in the plane, find the smallest circle enclosing them all.
- The optimal circle will be determined by at most 3 points, leading to an $O(n^4)$ “brute force” solution.
- With randomized incremental construction, however, we can solve this problem in only $O(n)$ expected time!

13

3 Simple Steps...

- Suppose by magic that we already know 2 points on the boundary of an optimal circle...
 - Now it's easy to compute the answer in $O(n)$ time:
 - Let p_1 and p_2 be the points we know.
 - Start with a circle C having p_1 and p_2 as endpoints of its diameter.
 - Let the remaining points $p_3 \dots p_n$ be arbitrarily ordered.
 - Process $p_3 \dots p_n$ in sequence, enlarging C when necessary so it remains an optimal circle for the set of points considered thus far.

14

3 Simple Steps...

- Suppose by magic that we already know 2 points on the boundary of an optimal circle...
 - Now it's easy to compute the answer in $O(n)$ time.
- Now suppose (also by magic) that we already know only 1 point on the boundary of an optimal circle...
 - We can still compute the optimal circle in $O(n)$ expected time...
 - Let p_1 be the point we know, and let $p_2 \dots p_n$ be randomly ordered. Start with a zero-area circle C centered at p_1 .
 - Process $p_2 \dots p_n$ in sequence, updating C as we go.

15

3 Simple Steps...

- Suppose by magic that we already know 2 points on the boundary of an optimal circle...
 - Now it's easy to compute the answer in $O(n)$ time.
- Now suppose (also by magic) that we already know only 1 point on the boundary of an optimal circle...
 - We can still compute the optimal circle in $O(n)$ expected time...
- Finally, suppose we know none of the points on the boundary of an optimal circle...

16


3 Simple Steps...

- Suppose by magic that we already know 2 points on the boundary of an optimal circle...
 - Now it's easy to compute the answer in $O(n)$ time.
- Now suppose (also by magic) that we already know only 1 point on the boundary of an optimal circle...
 - We can still compute the optimal circle in $O(n)$ expected time...
- Finally, suppose we know none of the points on the boundary of an optimal circle...
- Final running time: $O(d! n)$, where d is the dimensionality of our space.

17

The Convex Hull Problem

- The **convex hull** of a set of n points is the smallest convex polygon containing all n points:




- The convex hull problem:
 - **Input:** A list of n points $(x_1, y_1) \dots (x_n, y_n)$.
 - **Output:** An array or linked list specifying the points around the boundary of the hull in clockwise (or counterclockwise) order.

18

Farthest Pair of Points (2D)


- **Claim:** The farthest pair of points in a point set lies on the convex hull.
 - The two points will be on “opposite sides” of the hull.
- Once we’ve found the convex hull, we can therefore find the farthest pair of points in $O(n)$ time using the “rotating calipers” algorithm.
(we could call this a “sort and scan” approach...)



19

Collinear Points and Other Special Cases...

- Do all of these points belong to the convex hull, or just the endpoints?





- In general, computational geometry problems tend to be plagued by special cases like this.
 - This plus the danger of round-off errors can make it somewhat tricky to implement computational geometry algorithms correctly in practice!
- Typical assumption: no 3 of our points are collinear.

20

Primitive Operations

- Consider the following simple geometric questions:
 - Point P on line L?
 - Points P and Q on same side of line L?
 - Line segments S_1 and S_2 intersect?
 - Point P in angle formed by two rays R_1 and R_2 ?
 - Point P in convex polygon Q?
- A single “trick” makes all of these easy: look at the sign of the z component of the cross product!




$$(u \times v)_z = \begin{vmatrix} u_x & u_y \\ v_x & v_y \end{vmatrix}$$

21

Checking Convexity

- Using the cross product test on consecutive pairs of vectors v , we can test a polygon for convexity:




- As we walk around the polygon in a counter-clockwise direction, we should make only “left turns” (z component of cross product positive).


22

Convexifying a Polygon

- By splicing out points at which we turn right, we can “convexify” any polygon in $O(n)$ time.




- This gives a simple $O(n \log n)$ algorithm (known as the “package wrapping” algorithm) for convex hulls:
 - Pick a “reference” point q known to be on the hull (e.g, with minimum y coordinate).
 - Sort remaining points by angle from q .
 - Use this sorted ordering to build a polygon and then convexify it.



23


“QuickSort”



- How do we link two disjoint convex polygons in $O(n)$ time?

24


Linking Two Disjoint Polygons



- Splice polygons together into a non-convex polygon with a “doubled edge”, then convexify!

25


“Merge Sort”



- We can merge two (possibly overlapping) convex polygons in $O(n)$ time by sweeping monotonically around each one in order of angle.
 - Just like merging two sorted sequences, only slightly more geometric details to deal with along the way (further details omitted...)

26

“Merge Sort”




- We can merge two (possibly overlapping) convex polygons in $O(n)$ time by sweeping monotonically around each one in order of angle.
 - Just like merging two sorted sequences, only slightly more geometric details to deal with along the way (further details omitted...)

27

A Randomized Incremental Construction Convex Hull Algorithm


- We'll build up a convex hull by starting with 3 randomly-chosen points (a triangle) and adding the remaining points in random order, updating the hull when necessary.
- Maintain a point q inside the hull (the centroid of our initial triangle works fine).



28

Randomized Incremental Construction of a Convex Hull


- Maintain an ordered linked list of all points (+edges) in clockwise order around the current hull.
- For each point p not yet processed, maintain bi-directional pointers between p and the edge in our hull crossed by the ray qp (with this extra information, we can quickly test if p is inside or outside the current hull).



29

Randomized Incremental Construction of a Convex Hull


- When we add a new point p ,
 - Check in $O(1)$ time if p is inside or outside current hull.
 - If inside, nothing happens.
 - If outside, update the hull by removing necessary edges (shown in red below) and adding two new edges (the thick edges below); also **update the pointers of any affected yet-to-be-processed points**.



30

Running Time Analysis


- Add up running time per point p :
 - $O(1)$ total work during initialization.
 - $O(1)$ total work checking if p is inside or outside the current hull when p is finally selected.
 - $O(1)$ total work inserting p into the hull (if at all).
 - $O(1)$ total work later deleting p from the hull (if at all).
 - $O(\log n)$ expected total work for changing p 's “current hull edge” pointer during the course of the algorithm
(this is the only “expensive” part of the whole algorithm...)



31


Probability of Pointer Change: Thinking Backwards Again...

- Consider the point in time at which we have added k points to our instance (some on the current hull, some inside the hull).
- p : not-yet-added point with pointer to hull edge xy .
- Only deletion of x or y causes p 's pointer to change!
- We pick a point to delete at random from the set of all k outstanding points
- So w/ probability exactly $2/k$, p 's pointer changes when shrinking from a size- k point set to a size- $(k-1)$ point set.




Linearity of Expectation

- With probability exactly $2/k$, p 's pointer changes when shrinking from a size- k point set to a size- $(k-1)$ point set.
- X : Total number of changes to p 's pointer
- X_k : indicator r.v. telling us whether p 's pointer changes when adding the k th point.
- $E[X] = \sum_k E[X_k] = 2 \sum_k 1/k = O(\log n)$.
- And using randomized reduction, we can also argue that $E[X] = O(\log n)$.



Convex Hulls in Higher Dimensions


- Our randomized incremental construction algorithm extends readily to 3D, and keeps its $O(n \log n)$ expected running time!



- In $D > 3$ dimensions, the complexity of the hull can be as bad as $O(N^{\lfloor D/2 \rfloor})$. But we can compute convex hulls this quickly...

Running Time Analysis: 3D Convex Hull

- Add up running time per point p :
 - $O(1)$ total work during initialization.
 - $O(1)$ total work checking if p is inside or outside the current hull when p is finally selected.
 - ~~$O(1)$ total work inserting p into the hull (if at all).~~
 - ~~$O(1)$ total work later deleting p from the hull (if at all).~~
 - $O(\log n)$ expected total work for changing p 's “current hull edge” pointer during the course of the algorithm




35

Running Time Analysis: 3D Convex Hull


- Add up running time per point p :
 - $O(1)$ total work during initialization.
 - $O(1)$ total work checking if p is inside or outside the current hull when p is finally selected.
 - $O(1)$ expected work inserting p into the hull (if at all).
 - $O(1)$ expected work later deleting p from the hull (if at all).

(Euler's formula tells us that the average # of neighbors of a point in a polyhedron is $O(1)$, so the expected # of neighbors of a randomly-chosen point is also $O(1)$).




36


Some Prominent Problems in Computational Geometry...



Convex hull
Upper / lower hull



Lower / upper envelope



Voronoi diagram (solid)
Delaunay triangulation (dashed)

Geometric Duality

- In computational geometry, we often exploit geometric duality – exchanging the roles of points and lines.
 - Slope-intercept duality: point $(a,b) \leftrightarrow$ line $y = ax - b$.
 - Polar duality: point $(a,b) \leftrightarrow$ line $ax + by = 1$.
- The roles of points and lines reverse when we dualize:
 - Recall that two points determine a unique line, and likewise that two lines determine a unique point.
 - If L_1 and L_2 are two lines intersecting at point P , then $\text{dual}(L_1)$ and $\text{dual}(L_2)$ are two points that lie on the common line $\text{dual}(P)$.
 - If P_1 and P_2 are two points determining line L , then $\text{dual}(L)$ is the point of intersection between the lines $\text{dual}(P_1)$ and $\text{dual}(P_2)$.


Duality – Example

- Given N points in the 2D plane, find the two that determine a line with maximum slope.




Duality – Example

- Given N points in the 2D plane, find the two that determine a line with maximum slope.
- Apply slope-intercept duality: $(a,b) \leftrightarrow y = ax - b$
- Now we have a set of N lines, and we are looking for the two lines that determine a point with maximum coordinate!
- This point will occur between two lines with adjacent slopes (after sorting by slope); hence, in our original problem, the optimal line will occur between two points with adjacent x coordinates (after sorting by x coordinate).

Geometric Duality Between Upper/Lower Hulls and Lower/Upper Envelopes




Convex hull
Upper / lower hull




Lower / upper envelope

Planar Graph Duality Between Voronoi Diagrams and Delaunay Triangulations




- There are many ways (most rather complicated!) to compute both objects in 2D in $O(N \log N)$ time (and there is a matching lower bound)
- In $D > 2$ dimensions, the complexity of these objects can be as bad as $O(N^{\lfloor (D+1)/2 \rfloor})$.


Back to Our High-Level List of Problems...



Convex hull
Upper / lower hull



Lower / upper envelope




Voronoi diagram
Delaunay triangulation

We'll now discuss the *parabolic lifting* transformation, which allows us to transform between Voronoi/Delaunay problems in D dimensions and hull/envelope problems in $D+1$ dimensions!...

Parabolic Lifting

- Map each of our D -dimensional input points up to a $(D+1)$ -dimensional point by lifting it onto a paraboloid ($z = x^2 + y^2$ in 2D):



- Then draw tangent planes to the paraboloid at each of the N lifted points...

Parabolic Lifting

- By projecting the upper envelope of our tangent planes back down onto D dimensions, we get the Voronoi diagram of our original point set!
- By projecting the lower hull of our lifted points back down to D dimensions, we get the Delaunay triangulation of the original point set!

