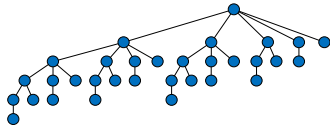




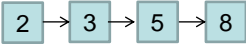
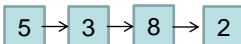
Lecture 16. Hashing and its Applications

CpSc 8400: Algorithms and Data Structures
Brian C. Dean



School of Computing
Clemson University
Spring, 2016


Dictionary/Set Data Structures

Insert	Remove	Find	
$O(n)$	$O(n)$	$O(\log n)$	Sorted array 
$O(1)$	$O(1)$, post-find	$O(n)$	Unsorted array 
$O(1)$, post-find	$O(1)$, post-find	$O(n)$	Sorted (doubly) linked list 
$O(1)$	$O(1)$, post-find	$O(n)$	Unsorted (doubly) linked list 
$O(\log n)$	$O(\log n)$	$O(\log n)$	Balanced BST, skip list, B-tree
$O(1)$ amortized	$O(1)$ amortized	$O(1)$ expected	Universal hash tables (today's topic)


First Attempt: A Direct Access Table

- Maintain a large array of bools
- Presence of key k in structure means $A[k]$ is true.

Contents of set:



Representation as array of bools:




- Insert, remove, and find all run in $O(1)$ time!
- Serious drawback: **space usage** (and therefore also initialization time).

3


Hash Tables

- Store elements in an array.
- Key k stored in position $h(k)$
- $h()$ is known as a “hash function” – it maps keys down to the range of indices in our array.
- Example: $h(k) = (2971k + 101923) \% 10$.

Contents of set:



Hash table:




$h(61) = 4$

4


Collision Resolution

- Probing: store elements directly in table; careful when removing elements.
- Chaining: store linked list of colliding elements off each table cell.



Another Way to Resolve Collisions: Cuckoo Hashing

- Key k stored at either position $h_1(k)$ and $h_2(k)$. (make sure h_1 and h_2 map to different positions!)
- Find, remove now trivial, and $O(1)$ time.
- Insert may lead to a chain of “evictions”.
 - If this goes on for too long (usually $O(\log n)$ steps), give up and rebuild the entire table with two new hash functions h_1 and h_2 .
- When table gets too full, rebuild at 2x size, just as before.



Why “Hashing”...?

Hash (n) A dish of cooked meat cut into small pieces and recooked, usually with potatoes.



7

Why “Hashing”...?

Hash (n) A dish of cooked meat cut into small pieces and recooked, usually with potatoes.



Hash (v) To jumble or mix up.

8

Choosing a Good Hash Function In Practice

- We want $h()$ to behave somewhat “randomly”.
 - Good Examples:
 - $h(k) = k \bmod m$.
 - $h(k) = \lfloor mk / C \rfloor$.
 - $h(k) = (ak) \bmod m$.
 - $h(k) = (ak + b) \bmod m$.
- m = table size
 C = max possible
key value
- } a and b should be chosen in a somewhat “arbitrary” fashion (primes often used).
- Be sure to fully utilize “all the bits” in a key. For example, $h(k) = k \bmod 256$ is a somewhat poor hash function if k is a 32-bit IP address.
- Be sure to use the entire hash table; e.g., make sure $h(k)$ isn’t always even.

9

Amortized Rebuilding

- How do we pick the initial size of our table?
- We ideally want to keep $m = \Theta(n)$, where m = table size and n = # of elements stored in table.
- To do this, double table size and re-hash when table becomes too full, and halve table size and re-hash when table becomes too empty.
- Since it takes only linear time to re-build the table, this makes insert and delete take + $O(1)$ extra amortized time.

Hashing Arrays

- Any complicated object can be “serialized” and represented by an array of small integers.
E.g., “Cpsc” \rightarrow ‘C’, ‘p’, ‘s’, ‘c’ \rightarrow 67, 112, 115, 99.
- So how do we hash an array $A[0 \dots N - 1]$ to get an output value in the range $0 \dots M - 1$?

11

Hashing Arrays

- Any complicated object can be “serialized” and represented by an array of small integers.
E.g., “Cpsc” \rightarrow ‘C’, ‘p’, ‘s’, ‘c’ \rightarrow 67, 112, 115, 99.
- So how do we hash an array $A[0 \dots N - 1]$ to get an output value in the range $0 \dots M - 1$?
- How about something like this:
$$h(A[]) = (A[0] + A[1] + \dots + A[N - 1]) \bmod M$$

12

Polynomial Hash Functions

- Think of $A[]$ as the coefficients of a polynomial:

$$p_A(x) = A[0] + A[1]x + A[2]x^2 + \dots + A[N-1]x^{N-1}$$

- To hash $A[]$, evaluate $p(x) \bmod M$ at a randomly-chosen point x . How do we do this quickly?
- If M is prime, then the probability two different arrays A and B collide at most $(N-1) / M$:
 - If A and B collide, then x is a root of $p_A(x) - p_B(x) \pmod{M}$.
 - A polynomial of degree $N-1$ can have at most $N-1$ roots (true in any algebraic field, such as arithmetic over complex numbers or over integers modulo a prime)

13

The Case Against Deterministic Hash Functions

- Any deterministic hash function has a “bad” set of input keys...
- If we are hashing n keys in the range $0 \dots C-1$ down to a table $A[0 \dots m-1]$ and $C \geq m(n-1) + 1$, then some set of $\geq n$ different keys will be mapped to the same table cell.
- This reduces our hash table to nothing more than a fancy linked list, so *find* takes $\Theta(n)$ time!
- So if we want a good worst-case guarantee for *find*, we cannot use a deterministic hash function. We **must** use randomness in some way.

14

Why Not Use a Completely Random Hash Function?

- Suppose we choose a hash function $h(k)$ that maps every key $k \in \{0, \dots, C - 1\}$ to a completely random location in $\{0, \dots, m - 1\}$.
- This gives us $\mathbf{E}[L] = n / m$, so *find* does indeed run in $O(1 + n / m) = O(1)$ time.
 - Please always assume (from here on) that we use dynamic resizing to maintain $m = \Theta(n)$, so $n / m = O(1)$.
 - How do we show that $\mathbf{E}[L] = n / m$? Linearity of expectation!
- Fatal flaw: ...?

15

Why Not Use a Completely Random Hash Function?

- Suppose we choose a hash function $h(k)$ that maps every key $k \in \{0, \dots, C - 1\}$ to a completely random location in $\{0, \dots, m - 1\}$.
- This gives us $\mathbf{E}[L] = n / m$, so *find* does indeed run in $O(1 + n / m) = O(1)$ time.
 - Please always assume (from here on) that we use dynamic resizing to maintain $m = \Theta(n)$, so $n / m = O(1)$.
 - How do we show that $\mathbf{E}[L] = n / m$? Linearity of expectation!
- Fatal flaw: requires $\Theta(C)$ space to store the hash function, same as the direct access table.

16

Universal Hashing

- A (randomized) hash function is **universal** if the probability (over random parameters in the function) of two different keys colliding is $O(1 / m)$.
- Example of a universal hash function (several others are discussed in the book):

$$h(k) = [(ak + b) \bmod p] \bmod m, \text{ with } \begin{cases} p: \text{any prime number } \geq C. \\ a: \text{random integer in } \{1, \dots, p-1\} \\ b: \text{random integer in } \{0, \dots, p\}. \end{cases}$$

- With a universal hash function, *find* runs in $O(1)$ expected time with chaining!
 - Easy proof using linearity of expectation...

17

$O(1)$ Expected Running Time for Find With Universal Hashing

- Let T denote the running time of an unsuccessful call to $\text{find}(k)$.
 - Clearly, $\mathbf{E}[T] \geq$ the expected time required for a successful find operation, so our analysis also provides a bound on the expected time of a successful find.
- Compute $\mathbf{E}[T]$ using linearity of expectation:
 - Let $k_1 \dots k_n$ denote the keys stored in our hash table.
 - Write $T = X_1 + X_2 + \dots + X_n$.
 - X_j : indicator random variable taking value 1 if $h(k_j) = h(k)$.
 - Since $k \neq k_j$, $\mathbf{E}[X_j] \leq 1 / m$ (using universal hashing).
 - So $\mathbf{E}[T] = \mathbf{E}[X_1] + \dots + \mathbf{E}[X_n] = n \mathbf{E}[X_j] \leq n / m = O(1)$.

18

Example Applications

- Consider the following problems:
 - **Element Uniqueness**: Given n numbers $A_1 \dots A_n$, are all of these distinct, or are two of them equal?
 - **Set Intersection / Union / Difference**: Given two sets A and B (specified by unsorted arrays) containing n elements in total, output an array containing $A \cap B$, $A \cup B$, or $A \setminus B$.
- All of these problems have $\Omega(n \log n)$ worst-case lower bounds in the comparison model.
- However, we can solve them in $O(n)$ expected time with universal hashing.

19

Levels of Independence Needed for $O(1)$ Expected Performance

- “Strongly” universal hashing functions are not fully random, but provide a weaker guarantee of pairwise independence.
 - If we only look at 2 keys at a time, these are hashed in a completely random fashion, just as with a fully random function (but not for 3 or more keys...)
- Higher levels of independence are possible using higher-degree polynomial hash functions. For example, this hash function is 4-wise independent:
$$h(k) = [(ak^3 + bk^2 + ck + d) \bmod p] \bmod m$$
- Recent result [Pagh et al. '11, Patrascu-Thorup '10]: 5-wise independence necessary and sufficient for $O(1)$ expected performance with linear probing.

20

Hashing: Much More than Just a Good Set Data Structure...

- The general idea of mapping a large, complicated object down to a simpler object appears in many areas of computing...
- Example: by hashing a file down to a small integer “fingerprint”, we can:
 - Compare (approximately) two files extremely quickly.
 - Test if two files at different locations (e.g., original, backup) are identical with a minimum amount of communication.
 - Detect tampering in critical system files.
 - Ensure integrity of the file when transferred over a noisy channel, by appending the fingerprint as a checksum.

21

Hashing in Security

- We can detect tampering in data (a file, or a message) if we also store a hash of the data.
 - Particularly if we use a cryptographic hash function (e.g., MD5, SHA1), which has specifically been designed to be hard to invert.
 - Difficult to change the data while keeping the same hash: by the birthday paradox, it takes $\sim 2^{64}$ guesses to find a single colliding pair of keys if our hash output is a 128-bit integer.
- Store passwords on a computer system as hashes instead of “in the clear” – still allows you to log in!

22