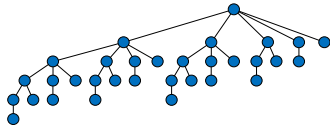


Lecture 20. Heuristics and Iterative Refinement


CpSc 8400: Algorithms and Data Structures
Brian C. Dean




School of Computing
Clemson University
Fall, 2016

Discrete Optimization

- We've already seen how to develop greedy algorithms and to use dynamic programming.
- For harder problems:
 - Approximation algorithms try to generate a provably-close-to-optimal solution in polynomial time.
 - Heuristics try to do as well as possible (either in terms of solution quality or running time) on real-world inputs.



Clustering




Rank aggregation



The traveling salesman problem (TSP)

Continuous Optimization

- Solving and optimizing not much different: for example, to solve $f(x) = 0$, minimize $|f(x)|^2$.
 - Often want to solve an entire system of equations (e.g., with finite element modeling).
- Regression is the process of fitting parameters to a model to minimize its error when fitting to data.




Iterative Refinement


- Simple yet powerful idea:
 - Start with some arbitrary feasible solution (or better yet, a solution obtained via some other heuristic)
 - As long as we can improve it, keep making it better (e.g., search a small “neighborhood” of solutions similar to x , moving to better one if found)
 - Simple example: bubble sort.
- Getting stuck in local minima can be a problem.
- Many approaches refine a population of solutions, rather than a single solution.

Local Search – Neighborhood Examples

Traveling salesman:



Rank aggregation:




Neighborhood size? (large vs. small)

Movement strategy? (immediate vs. after searching entire neighborhood)

5


Unconstrained Optimization of $F(x)$ Based On Iterative Refinement

- To refine a guess x , locally approximate F at x using either a linear function or a quadratic.



Gradient Descent (if derivatives easy to compute)

- Careful with step size
- Many fast “low quality” steps




Newton's Method (if second derivatives easy to compute)

- Fewer slow “high quality” steps (in N dimensions, each step involves solving an N -variable linear system).

- Convex functions nice, since local optimal point is also a global optimum point.


Unconstrained Optimization of $F(x)$ Based On Iterative Refinement

- To refine a guess x , locally approximate F at x using either a linear function or a quadratic.



Gradient Descent (if derivatives easy to compute)

- Careful with step size
- Many fast “low quality” steps




Newton's Method (if second derivatives easy to compute)

- Fewer slow “high quality” steps (in N dimensions, each step involves solving an N -variable linear system).

What if computing derivatives at all isn't easy...?

If Derivatives are Hard to Compute... “Ameoba Search”

- Goal: Minimize $F(x)$ without computing any derivatives...



- In higher dimensions, this leads to an idea known as downhill simplex search, the Nelder-Mead algorithm, or “amoeba” search.

Simulated Annealing

- Like neighborhood search, only you are allowed to move to worse solutions (with a goal of escaping local minima).
- Probability of moving to a worse solution proportional to how much worse it is, and this probably decreases over time as well according to a “cooling” schedule.
- Think of this as a “random walk” through different solutions, where it’s much more likely to move in an improving direction.

Population-Based Search: Genetic Algorithms



- Start with a **population** of initial solutions.
- Each member of next generation obtained by:
 - Mutation of good solution from previous generation,
 - Result of “mating” together two good solutions from previous generation, or
 - Direct copy of best solution from previous generation.
- Solutions from previous generation chosen with probability related to their **fitness**.
 - Sometimes called “roulette wheel” selecton.

Genetic Algorithms: “Crossover” for Combining Solutions...


$x :$		$x' :$
1	1	1
0	1	1
0	1	1
1	0	0
0	1	1
1	1	1
1	1	0
0	0	1

Genetic Algorithms: “Crossover” for Combining Solutions...

$x :$		$x' :$
1	1	1
0	1	1
0	1	1
1	0	0
0	1	1
1	1	1
1	1	0
0	0	1




Genetic Algorithms: “Crossover” for Combining Solutions...



Physical and Biologically Inspired Computing Paradigms

- Simulated Annealing
- Genetic/Evolutionary Algorithms
- Bio-Inspired Optimization Methods
- Neural Networks
- Artificial Immune Systems



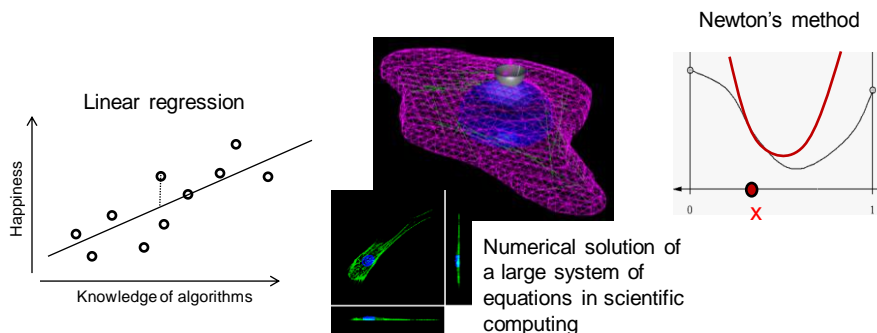
Linear Systems

- N unknown variables, N equations
(what if # of variables is more or less than the # of equations?)
- Example: $x + y = 3$
 $2x - y = 3$
- Example: $x_1 + 2x_2 + 3x_3 = 10$
 $4x_1 + 5x_2 + 6x_3 = 11$
 $7x_1 + 8x_2 + 9x_3 = 12$
- We often write these using matrix-vector notation:

$$\begin{matrix} \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} & \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} & = & \begin{bmatrix} 10 \\ 11 \\ 12 \end{bmatrix} \\ A & x & & b \end{matrix}$$

Linear Systems are Everywhere

- Least squares regression or solving a system of equations usually involves minimizing a quadratic function giving squared error.
- To minimize a quadratic, set its gradient to zero – this gives a linear system!



Solving Linear Systems with Gaussian Elimination

- Consider solving an $N \times N$ linear system $Ax = b$.
- Example:
$$\begin{aligned}x_1 + 2x_2 + 3x_3 &= 10 \\4x_1 + 5x_2 + 6x_3 &= 11 \\7x_1 + 8x_2 + 9x_3 &= 12\end{aligned}$$
- Solve first equation for x_1 in terms of other variables: $x_1 = 10 - 2x_2 - 3x_3$.
- Plug this in to eliminate x_1 from equations 2...N, giving a system of $N - 1$ variables, $N - 1$ equations.
- Eventually, we get a single equation in x_N . Solve it, then work backwards to get values of $x_{N-1} \dots x_1$.
- Total running time: $O(N^3)$

Solving Linear Systems with Gaussian Elimination

- The “standard” method for solving a linear system via Gaussian elimination takes $O(N^3)$ time.
- With a bit of cleverness, we can solve a linear system in the same amount of time it takes to multiply two $N \times N$ matrices (also $O(N^3)$ time using the most straightforward approach).
- However, matrix multiplication can be done a bit faster...

Solving Linear Systems with Gaussian Elimination

- The “standard” method for solving a linear system via Gaussian elimination takes $O(N^3)$ time.

- Matrix Multiplication...

Best known lower bound: $\Omega(n^2)$ (trivial)

Strassen 1969: $O(n^{2.81})$

- Coppersmith and Winograd 1990: $O(n^{2.376})$
(However, it's quite complicated and not very practical)

Stothers 2011: $O(n^{2.374})$

Williams 2012: $O(n^{2.3729})$

Le Gall 2014: $O(n^{2.37287})$

Solving Linear Systems (Potentially Faster) with Iterative Refinement

- Consider solving an $N \times N$ linear system $Ax = b$.

- Example: $x_1 + 2x_2 + 3x_3 = 10$

$$4x_1 + 5x_2 + 6x_3 = 11$$

$$7x_1 + 8x_2 + 9x_3 = 12$$

- A much faster solution in practice uses iterative refinement:

- Guess an initial solution $x = (x_1 \dots x_N)$
- Solve equation 1 for x_1 to update x_1 .
- Solve equation 2 for x_2 to update x_2 .
- Etc...
- Repeating entire process until convergence.

