


Lecture 22. Shortest Paths

CpSc 8400: Algorithms and Data Structures
Brian C. Dean





School of Computing
Clemson University
Spring, 2016

Types of Shortest Path Problems

- **Single-source single-destination:** find the shortest path from s to t in a (directed) graph.
- **Single-source:** find the shortest path from s to every other node in a (directed) graph.
- **All-pairs:** find the shortest path from every node to every other node.

Single-Source Shortest Paths



Input:

Directed graph with costs/weights/lengths on its edges.

Output:

- $c[j]$: cost of shortest $s \rightarrow j$ path
- $pred[j]$: previous node (before j) on a shortest $s \rightarrow j$ path.

3

Single-Source Shortest Paths: Algorithms

- Easy cases that we can solve in linear time:
 - Unweighted graphs: use **breadth-first search**.
 - Directed acyclic graphs (DAGs): use **dynamic programming**.
- If edge costs are nonnegative, we'll use **Dijkstra's algorithm**.
- If some edge costs are negative, we use the **Bellman-Ford algorithm**.

4

Breadth-First Search

- Depth-first search dives as deeply as possible into a graph until it can go no further, then it backs up and tries to branch.
- In contrast, a **breadth-first search** (BFS) starting at some source node s will visit s , then all nodes 1 hop away from s , then all nodes 2 hops away from s , etc.
- Like DFS, we can also use BFS to find connected components or answer “find a path from i to j ” queries.
 - BFS finds a path from i to j having fewest edges.

5

Breadth-First Search

```
BFS(s):  
  For all nodes i:  
    pred[i] = null  
    dist[i] = Infinity  
  Q = {s}  
  dist[s] = 0  
  While Q is nonempty:  
    i = next node in Q  
    For all nodes j such that (i,j) is an edge:  
      If dist[j] = Infinity,  
        pred[j] = i  
        dist[j] = dist[i] + 1  
        Append j to the end of Q
```

- If Q is implemented as a stack rather than as a FIFO queue, we get DFS rather than BFS!

6


Negative-Cost Cycles

- If negative-cost edges exist:
 - We need to use the $O(mn)$ Bellman-Ford algorithm.
 - We can only find shortest paths if there are no negative-cost directed cycles (Bellman-Ford can at least detect whether negative cycles exist).
 - If our graph is undirected, we typically assume it has no negative-cost edges to begin with.
- Longest path problems convert to shortest path problems by negating edge lengths, so positive-cost cycles are bad for these problems
 - So longest path problems are rarely tractable, since most graphs have positive-cost cycles.
 - Except DAGs, that is...

7

The Triangle Inequality

- Shortest path costs satisfy the triangle inequality:
 $c[i] + \text{cost}(i, j) \geq c[j]$ for every edge (i, j) .



- **Shortest path optimality conditions:** If we have a set of valid shortest path cost labels $c[j]$ that satisfy the triangle inequality for every edge, then these costs must be optimal.

8

Fixing the Triangle Inequality

- If we ever notice an edge (i, j) that violates the triangle inequality:

$$c[i] + \text{cost}(i, j) < c[j]$$

then we ought to fix things by reducing $c[j]$:

```
Tighten(i, j):  (also called Relax(i, j))
  If  $c[i] + \text{cost}(i, j) < c[j]$ ,
     $c[j] = c[i] + \text{cost}(i, j)$ 
     $\text{pred}[j] = i$ 
```

- Most shortest path algorithms start with $c[s] = 0$ and $c[j \neq s] = +\infty$, and repeatedly tighten (relax) edges violating the Δ inequality.
 - This guarantees a set of valid cost labels $c[j]$ that will be optimal when we terminate.

9


Dijkstra's Algorithm

```
Dijkstra:
   $c[s] = 0$ ,  $c[j \neq s] = +\infty$ 
  Build a priority queue  $Q$  on the  $c[j]$ 's
  While  $Q$  is nonempty:
    Remove from  $Q$  the node  $i$  having minimum label  $c[i]$ 
      (at this point we know  $c[i]$  is correct)
    Tighten all edges emanating from  $i$ 
```

- Only works if edge costs are nonnegative.
- Running time depends on priority queue:
 - n inserts (during initialization)
 - n remove-mins
 - at most m decrease-keys
- Using a binary heap, all these operations run in $O(\log n)$ time, so total is $O(m \log n)$.

10


Dijkstra's Algorithm – Example



Initialization.

11


Dijkstra's Algorithm – Example



Examine the source node s , so s becomes “permanently” labeled.
Tighten all edges leaving s .

12


Dijkstra's Algorithm – Example



Examine node 3, which becomes permanently labeled.
Tighten all edges leaving node 3.

13


Dijkstra's Algorithm – Example



Examine node 1, which becomes permanently labeled.
Tighten all edges leaving node 1.

14


Dijkstra's Algorithm – Example



Examine node 4, which becomes permanently labeled.
Tighten all edges leaving node 4.

15


Dijkstra's Algorithm – Example



Examine node 5, which becomes permanently labeled.
Tighten all edges leaving node 5 (there aren't any).

16

Dijkstra's Algorithm – Example



Examine node 2, which becomes permanently labeled.
Algorithm terminates!

17

Dijkstra's Algorithm – Running Time

	Remove-Min $O(n)$ times	Decrease Key $O(m)$ times	Total Runtime
Unsorted Array	$O(n)$	$O(1)$	$O(n^2)$
Binary Heap	$O(\log n)$	$O(\log n)$	$O(m \log n)$
Fibonacci Heap	$O(\log n)$	$O(1)$ amortized	$O(m + n \log n)$

- The Fibonacci heap bound is always at least as good as both other bounds in any graph!
- In the RAM model (small integer edge costs), we can improve performance by using special RAM priority queues.

18

Dijkstra's Algorithm – Correctness

- Just like BFS, Dijkstra visits all nodes in increasing order of distance from the source, and visits each node exactly once.
- **Claim:** when we visit node i , $c[i]$ will be the correct shortest path cost from s to i .
- Think of Dijkstra as maintaining two sets:
 - S: visited nodes that are “permanently” labeled.
 - T: remaining nodes that are temporarily labeled (whose cost labels are still upper bounds)
- Every iteration takes the node from T with minimum label and adds it to S.

19

The Bellman-Ford Algorithm

```
Bellman-Ford:
  c[s] = 0, c[j ≠ s] = +∞
  Repeat n-1 times:
    Tighten every edge in the graph
  If c[i] + cost(i, j) < c[j] for any edge (i, j)
    Output "negative cycle detected!"
```

- $O(mn)$ running time.
- Detects negative cycles.
- Easy to analyze using induction: after k iterations of the main loop, each label $c[j]$ reflects the cost of the shortest k -hop path from s to j .
- Can also interpret as a dynamic programming algorithm, or equivalently as our algorithm for finding shortest paths in a DAG with n “levels”.

20

All-Pairs Shortest Path Algorithms

- Typically we think of this problem as taking an $n \times n$ matrix of edge costs as input and producing an $n \times n$ matrix of shortest path costs as output.
- Related problem: find the **diameter** of a graph.
- Can solve the problem using n invocations of a single-source algorithm.
 - Nonnegative edge costs: $n \times \text{Dijkstra} = O(mn + n^2 \log n)$.
 - Negative edge costs: $n \times \text{Bellman-Ford} = O(mn^2)$.
(this can be improved to $O(mn + n^2 \log n)$ by using Johnson's algorithm: run Bellman-Ford once, then "reweight" the edges of the graph with nonnegative costs, and run Dijkstra n times...).

21

Johnson's Algorithm

- Add a dummy node s to the graph with a zero-cost edge to every other node.
- Use Bellman-Ford to find the shortest path cost $c[i]$ from s to every other node i .
- These costs $c[i]$ satisfy the triangle inequality:
 $c[i] + \text{cost}(i, j) \geq c[j]$ for every edge (i, j) .
- Define the **reduced cost** of (i, j) as:
 $\text{cost}'(i, j) = \text{cost}(i, j) + c[i] - c[j] \geq 0$.
- Now find shortest paths from every other source node using Dijkstra's algorithm!
- Total running time: $O(mn + n^2 \log n)$.

22

Johnson's Algorithm

- Define the **reduced cost** of (i, j) as:
 $\text{cost}'(i, j) = \text{cost}(i, j) + c[i] - c[j] \geq 0$.
- The reduced cost of an $i \rightarrow j$ path p is:
 $\text{cost}'(p) = \text{cost}(p) + c[i] - c[j]$.
- So when we convert to reduced costs, this just offsets the cost of every $i \rightarrow j$ path by the same amount: $c[i] - c[j]$. The shortest $i \rightarrow j$ path remains the shortest, however!
- Reduced costs have many nice uses in shortest path calculations.

23

The Floyd-Warshall Algorithm

- Let $c(i, j)$ denote the cost of an edge from node i to node j .
 - If there is no edge, set $c(i, j) = \infty$.
 - Also, set $c(i, i) = 0$.
- After running the remarkably simple $O(n^3)$ Floyd-Warshall algorithm...

Floyd-Warshall:

For $k = 1$ to n

For $i = 1$ to n

For $j = 1$ to n

If $c[i, k] + c[k, j] < c[i, j]$, $c[i, j] = c[i, k] + c[k, j]$

- ... $c(i, j)$ now gives us the cost of a shortest path from i to j !

24

The Floyd-Warshall Algorithm


```
Floyd-Warshall:  
For k = 1 to n  
  For i = 1 to n  
    For j = 1 to n  
      If  $c[i,k] + c[k,j] < c[i,j]$ ,  $c[i,j] = c[i,k] + c[k,j]$ 
```

- This works even if there are negative-cost edges. We can detect a negative cycle if $c(i, i) < 0$ for any node i after termination.
- Correctness follows from analysis as a dynamic programming algorithm:
 - Let $c^k[i, j]$ denote the cost of a shortest $i \rightarrow j$ path that uses only nodes $1 \dots k$ as potential intermediate nodes along the path.
 - Initially, $c^0[i, j]$ = the cost of edge (i, j) .
 - Moreover, $c^n[i, j]$ = the cost of a shortest $i \rightarrow j$ path.
 - DP recursion: $c^k[i, j] = \min\{c^{k-1}[i, j], c^{k-1}[i, k] + c^{k-1}[k, j]\}$

25

Recall: Centrality

- What nodes are most “central” in a graph?



26

Centrality & Shortest Paths

- Many notions of centrality are based on shortest paths:
 - Eccentricity(x) = max dist(x,y) over other nodes y ; a graph “center” is a node with minimum eccentricity.
 - “Betweenness” centrality of an edge or node: fraction of all point-to-point shortest paths passing through that edge/node.
- Other prominent notions of centrality include:
 - Degree centrality: high degree treated as “more central”
 - Eigenvector centrality (centrality proportional to probability a node is visited by a random walk in steady state); used by Google Pagerank.



27

Formulating Shortest Path Problems: Change the Problem, not the Algorithm...

- In an automobile transportation network, suppose right turns can be made instantly, while left turns take 1 extra unit of time...

28

Formulating Shortest Path Problems: Change the Problem, not the Algorithm...

- In an automobile transportation network, suppose right turns can be made instantly, while left turns take 1 extra unit of time...
- Suppose we have N jobs to complete. Each job has a duration and possibly a set of prerequisite jobs. Working in parallel, how quickly can we finish all the jobs?

29

Formulating Shortest Path Problems: Change the Problem, not the Algorithm...

- In an automobile transportation network, suppose right turns can be made instantly, while left turns take 1 extra unit of time...
- Suppose we have N jobs to complete. Each job has a duration and possibly a set of prerequisite jobs. Working in parallel, how quickly can we finish all the jobs?
- Suppose some roads charge a \$1 toll. How do we find a shortest path if we can pay at most \$3 worth of tolls?

30

Formulating Shortest Path Problems: Change the Problem, not the Algorithm...

- In an automobile transportation network, suppose right turns can be made instantly, while left turns take 1 extra unit of time...
- Suppose we have N jobs to complete. Each job has a duration and possibly a set of prerequisite jobs. Working in parallel, how quickly can we finish all the jobs?
- Suppose some roads charge a \$1 toll. How do we find a shortest path if we can pay at most \$3 worth of tolls?
- What is the second-shortest s - t path?

31